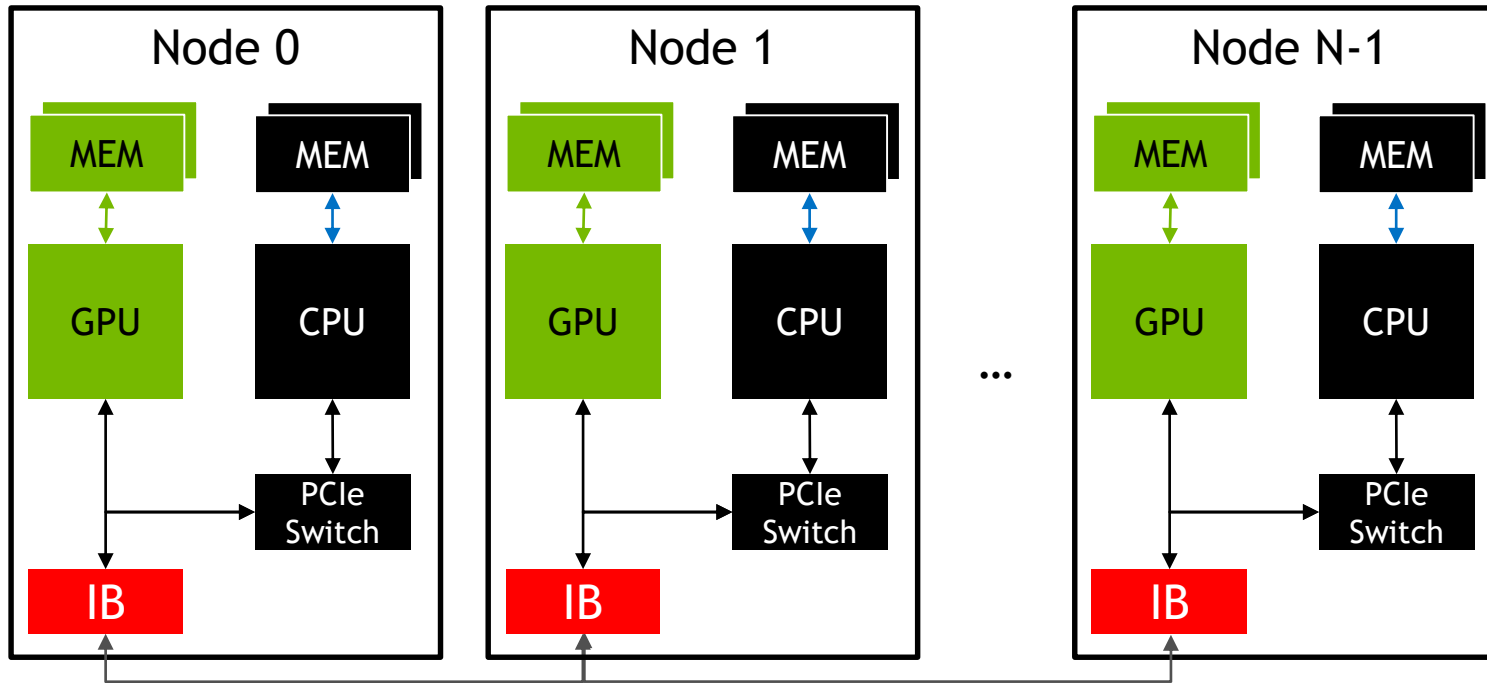




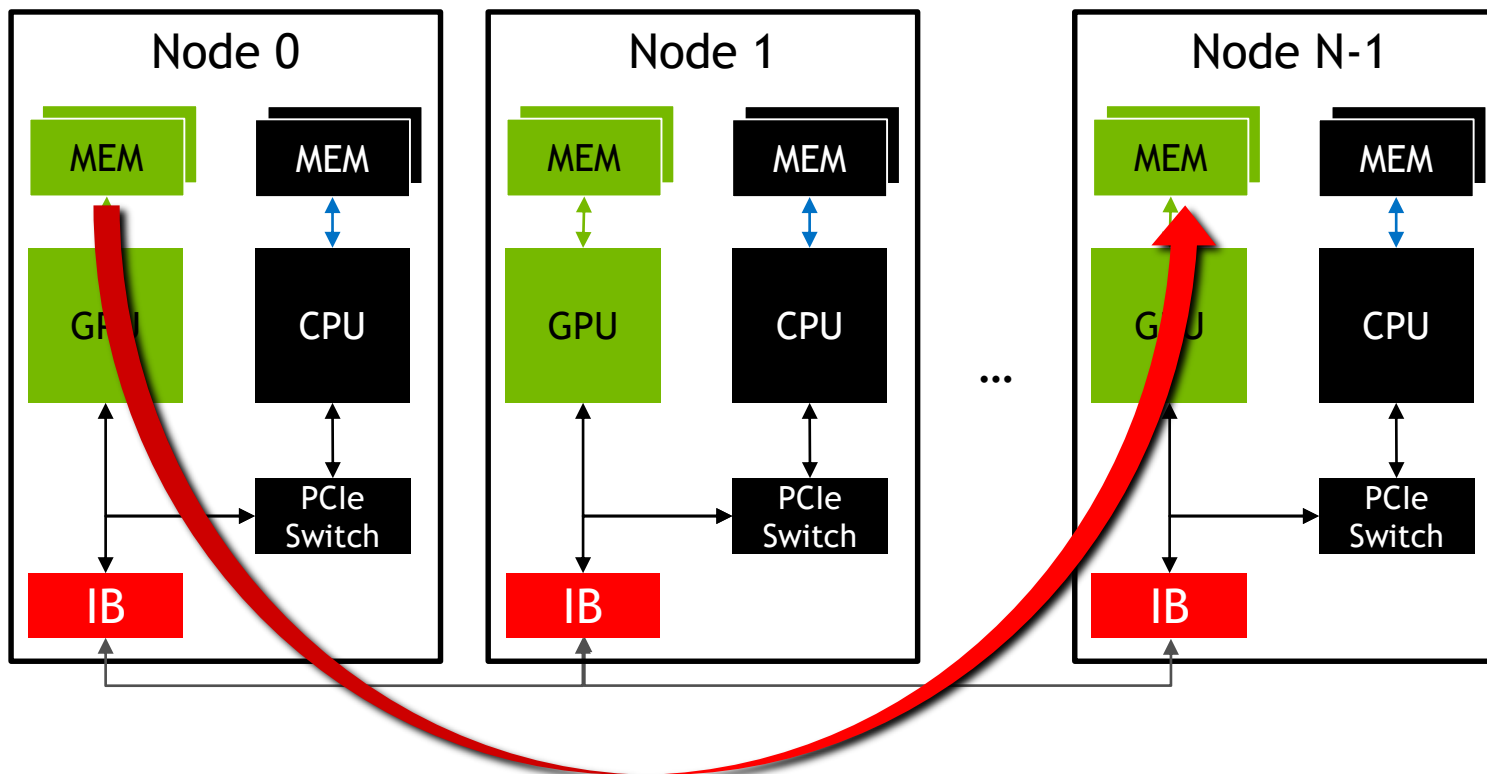
# Multi-GPU Programming with MPI

Cheng Yi, Senior Solution Architect, NVIDIA

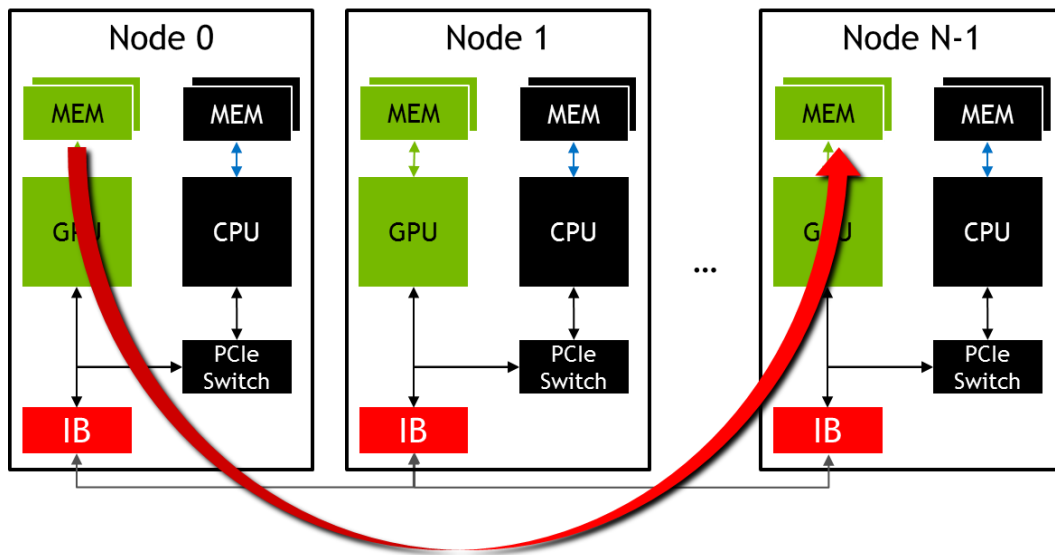
# MPI+CUDA



# MPI+CUDA



# MPI+CUDA



```
//MPI rank 0
MPI_Send(s_buf_d,size,MPI_CHAR,n-1>tag,MPI_COMM_WORLD);

//MPI rank n-1
MPI_Recv(r_buf_d,size,MPI_CHAR,0>tag,MPI_COMM_WORLD,&stat);
```

# YOU WILL LEARN

What MPI is

How to use MPI for inter GPU communication with CUDA and OpenACC

What CUDA-aware MPI is

What Multi Process Service is and how to use it

How to use NVIDIA tools in an MPI environment

How to hide MPI communication times

# MESSAGE PASSING INTERFACE - MPI

Standard to exchange data between processes via messages

Defines API to exchanges messages

Point to Point: e.g. `MPI_Send`, `MPI_Recv`

Collectives: e.g. `MPI_Reduce`

Multiple implementations (open source and commercial)

Bindings for C/C++, Fortran, Python, ...

E.g. MPICH, OpenMPI, MVAPICH, IBM Platform MPI, Cray MPT, ...

# MPI - SKELETON

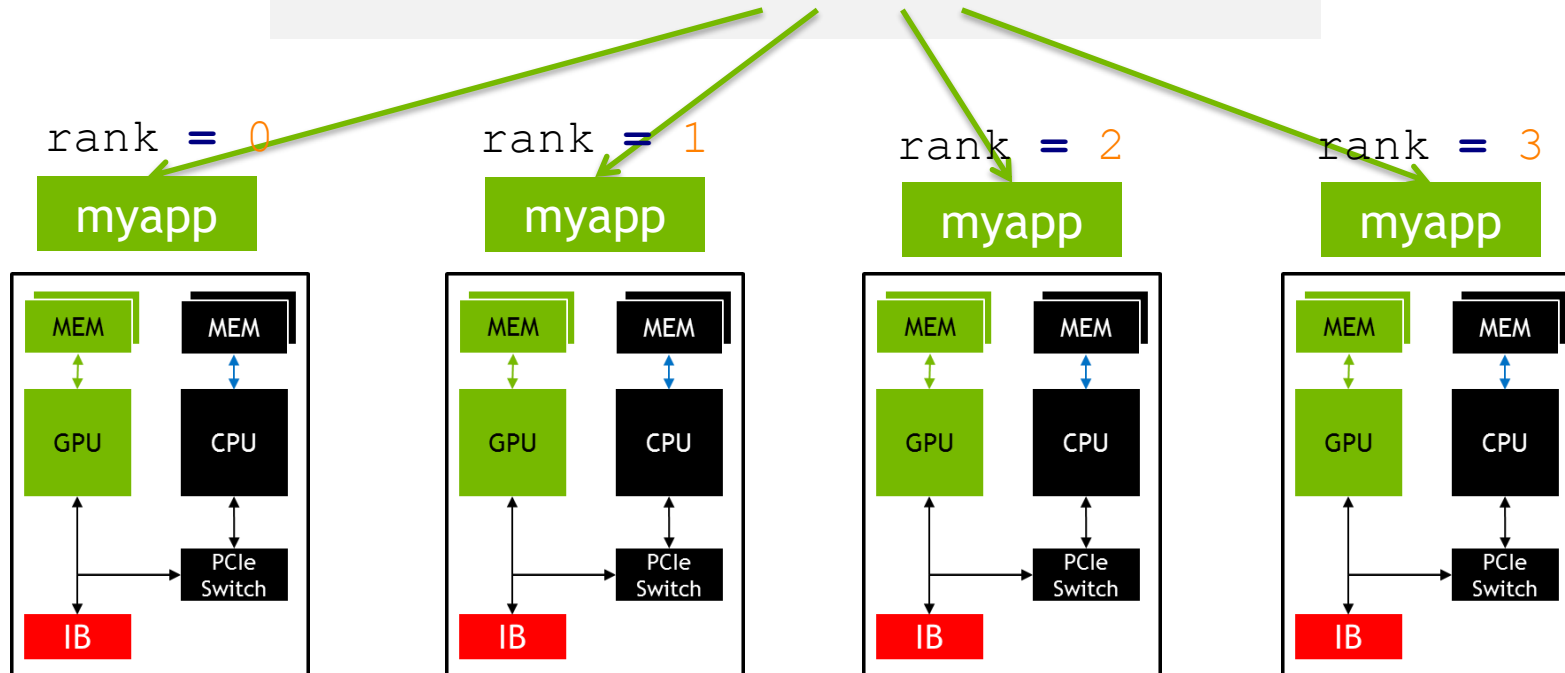
```
#include <mpi.h>

int main(int argc, char *argv[]) {
    int rank, size;
    /* Initialize the MPI library */
    MPI_Init(&argc, &argv);
    /* Determine the calling process rank and total number of ranks */
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    /* Call MPI routines like MPI_Send, MPI_Recv, ... */
    ...
    /* Shutdown MPI library */
    MPI_Finalize();
    return 0;
}
```

# MPI

## Compiling and Launching

```
$ mpicc -o myapp myapp.c  
$ mpirun -np 4 ./myapp <args>
```





# A SIMPLE EXAMPLE

# EXAMPLE: JACOBI SOLVER

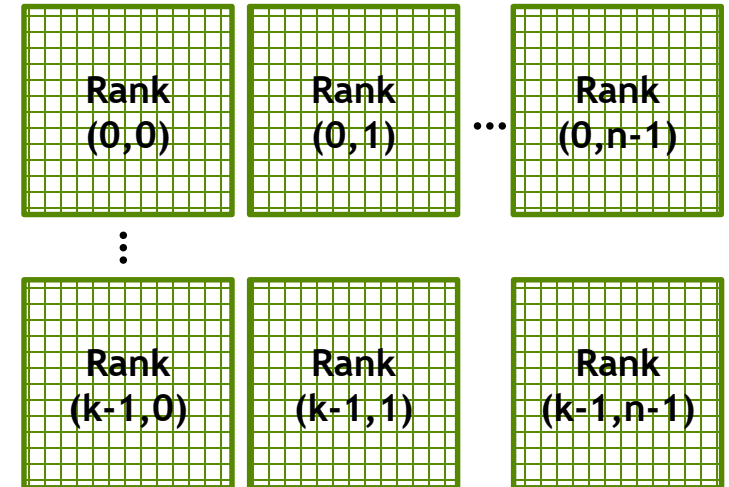

Solves the 2D-Laplace Equation on a rectangle

$$\Delta \mathbf{u}(x, y) = \mathbf{0} \quad \forall (x, y) \in \Omega \setminus \delta\Omega$$

Dirichlet boundary conditions (constant values on boundaries)

$$\mathbf{u}(x, y) = \mathbf{f}(x, y) \in \delta\Omega$$

2D domain decomposition with  $n \times k$  domains



# EXAMPLE: JACOBI SOLVER

## Single GPU

While not converged


Do Jacobi step:

```
for (int iy=1; iy < ny-1; ++iy)
for (int ix=1; ix < nx-1; ++ix)

    u_new[ix][iy] = 0.0f - 0.25f*( u[ix-1][iy] + u[ix+1][iy]
                                   + u[ix][iy-1] + u[ix][iy+1]);
```

Swap `u_new` and `u`

Next iteration



# EXAMPLE: JACOBI SOLVER

## Multi GPU

While not converged

Do Jacobi step:


```
for (int iy=1; iy < ny-1; ++iy)
for (int ix=1; ix < nx-1; ++ix)

    u_new[ix][iy] = 0.0f - 0.25f*( u[ix-1][iy] + u[ix+1][iy]
                                   + u[ix][iy-1] + u[ix][iy+1]);
```

Exchange halo with 2 4 neighbors

Swap `u_new` and `u`

Next iteration




# EXAMPLE JACOBI

## Top/Bottom Halo

```
MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,  
u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,  
MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```

```
MPI_Sendrecv(u_new+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,  
u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,  
MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```



# EXAMPLE JACOBI


## Top/Bottom Halo

OpenACC

```
#pragma acc host_data use_device ( u_new ) {  
MPI_Sendrecv( u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,  
              u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,  
              MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
MPI_Sendrecv( u_new+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,  
              u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,  
              MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```

CUDA

```
MPI_Sendrecv( u_new_d+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,  
              u_new_d+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,  
              MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
MPI_Sendrecv( u_new_d+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,  
              u_new_d+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,  
              MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```



# EXAMPLE: JACOBI



## Left/Right Halo

OpenACC

```
//right neighbor omitted
#pragma acc parallel loop present ( u_new, to_left )
for ( int i=0; i<n-2; ++i )
    to_left[i] = u_new[(i+1)*m+1];

#pragma acc host_data use_device ( from_right, to_left ) {
    MPI_Sendrecv( to_left, n-2, MPI_DOUBLE, l_nb, 0,
                  from_right, n-2, MPI_DOUBLE, r_nb, 0,
                  MPI_COMM_WORLD, MPI_STATUS_IGNORE );
}

#pragma acc parallel loop present ( u_new, from_right )
for ( int i=0; i<n-2; ++i )
    u_new[(m-1)+(i+1)*m] = from_right[i];
```



# EXAMPLE: JACOBI

## Left/Right Halo

CUDA



```
//right neighbor omitted
```

```
pack<<<gs,bs,0,s>>>(to_left_d, u_new_d, n, m);
```

```
cudaStreamSynchronize(s);
```

```
MPI_Sendrecv( to_left_d, n-2, MPI_DOUBLE, l_nb, 0,  
              from_right_d, n-2, MPI_DOUBLE, r_nb, 0,  
              MPI_COMM_WORLD, MPI_STATUS_IGNORE );
```

```
unpack<<<gs,bs,0,s>>>(u_new_d, from_right_d, n, m);
```





# LAUNCH MPI+CUDA/OPENACC PROGRAMS

Launch one process per GPU

MVAPICH: MV2\_USE\_CUDA

```
$ MV2_USE_CUDA=1 mpirun -np ${np} ./myapp <args>
```

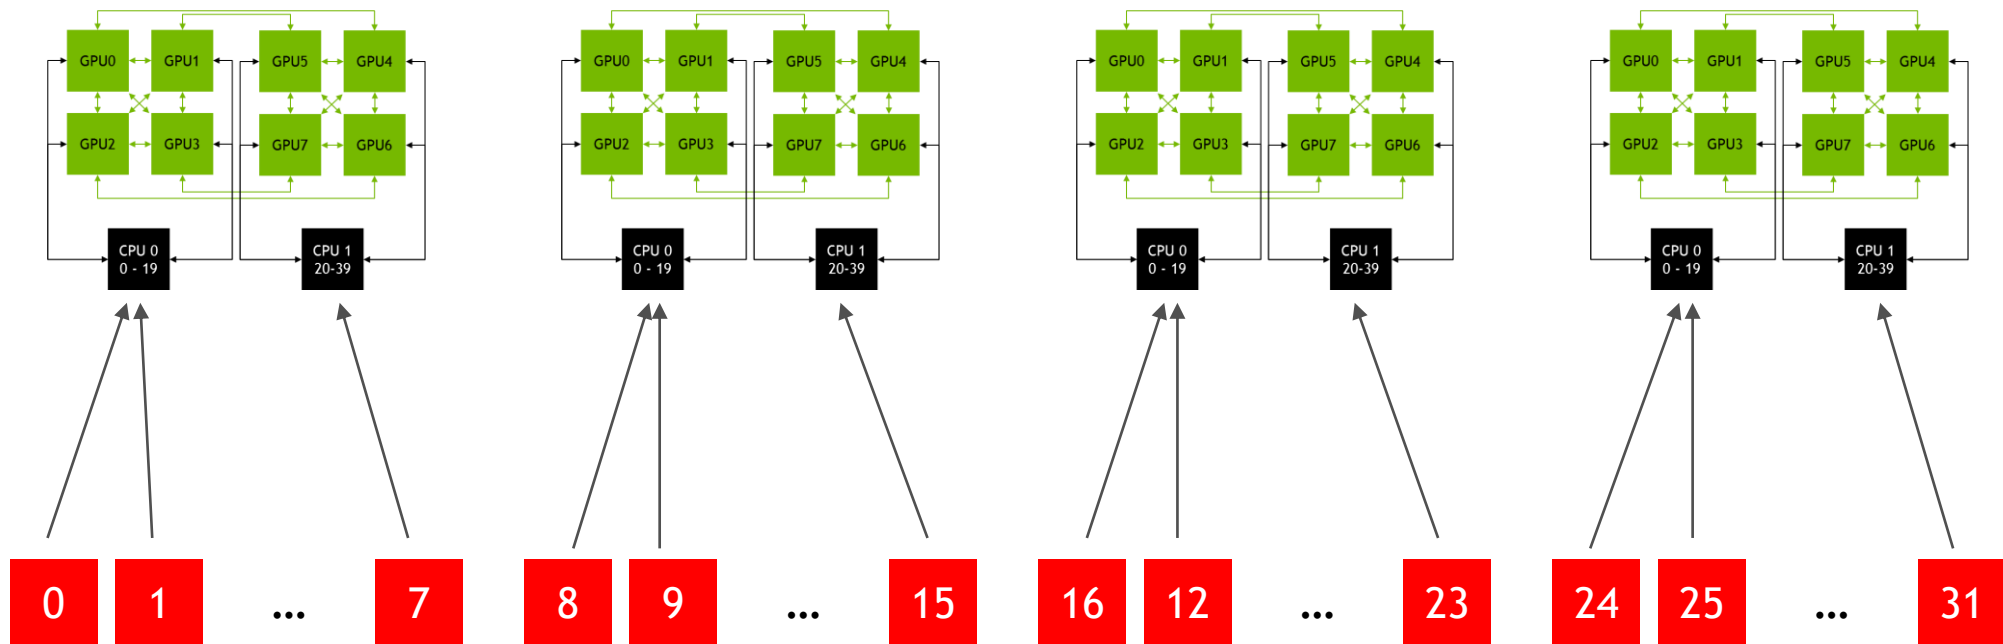
Open MPI: CUDA-aware features are enabled per default

Cray: MPICH\_RDMA\_ENABLED\_CUDA

IBM Spectrum MPI:

```
$ mpirun -gpu -np ${np} ./myapp <args>
```

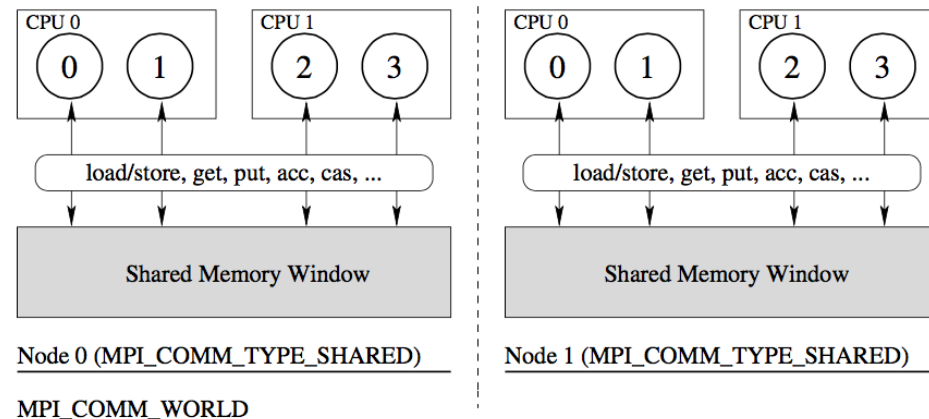
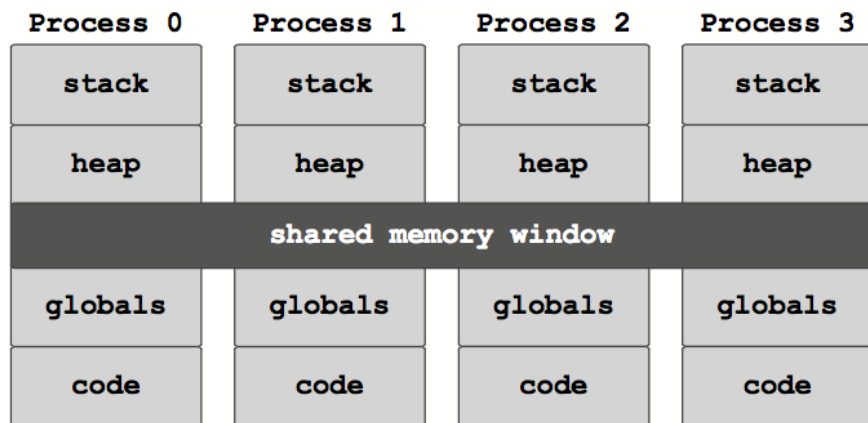
# HANDLING MULTIPLE MULTI GPU NODES



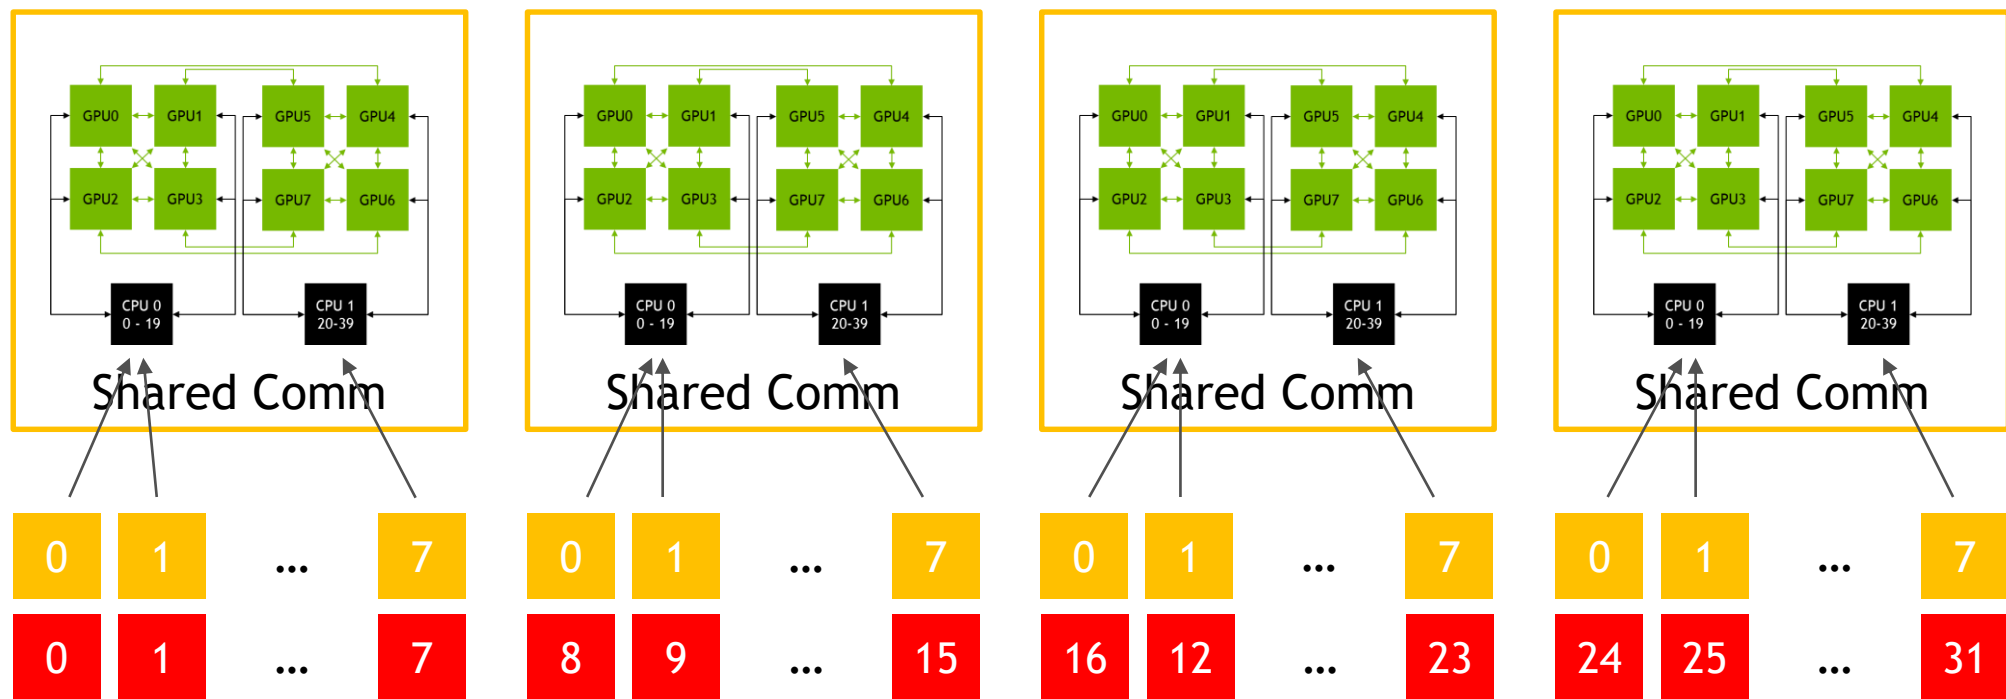
# HANDLING MULTIPLE MULTI GPU NODES

## How to determine the local rank? - MPI-3

```
MPI_Comm loc_comm;  
  
MPI_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED, rank, MPI_INFO_NULL, &loc_comm);  
  
int local_rank = -1;  
  
MPI_Comm_rank(loc_comm, &local_rank);  
  
MPI_Comm_free(&loc_comm);
```



# HANDLING MULTIPLE MULTI GPU NODES



# HANDLING MULTIPLE MULTI GPU NODES

## GPU-affinity

Use local rank:

```
int local_rank = -1;

MPI_Comm_rank(local_comm, &local_rank);

int num_devices = 0;

cudaGetDeviceCount(&num_devices);

cudaSetDevice(local_rank % num_devices);
```

# EXAMPLE JACOBI

## Top/Bottom Halo

without  
CUDA-aware  
MPI

OpenACC

```
#pragma acc update host(u_new[offset_first_row:m-2],u_new[offset_last_row:m-2])
MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,
             u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
MPI_Sendrecv(u_new+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,
             u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
#pragma acc update device(u_new[offset_top_boundary:m-2],u_new[offset_bottom_boundary:m-2])
```

CUDA

```
//send to bottom and receive from top top bottom omitted



cudaMemcpy( u_new+offset_first_row,
            u_new_d+offset_first_row, (m-2)*sizeof(double), cudaMemcpyDeviceToHost);
MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,
            u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,
            MPI_COMM_WORLD, MPI_STATUS_IGNORE);
cudaMemcpy( u_new_d+offset_bottom_boundary,
            u_new+offset_bottom_boundary, (m-2)*sizeof(double), cudaMemcpyDeviceToHost);
```

# THE DETAILS

# UNIFIED VIRTUAL ADDRESSING


*No UVA: Separate Address Spaces*

*UVA: Single Address Space*





# UNIFIED VIRTUAL ADDRESSING



One address space for all CPU and GPU memory

Determine physical memory location from a pointer value


Enable libraries to simplify their interfaces (e.g. MPI and cudaMemcpy)

Supported on devices with compute capability 2.0+ for

64-bit applications on Linux and Windows (+TCC)


# NVIDIA GPUDIRECT™

## Peer to Peer Transfers



# NVIDIA GPUDIRECT™

Support for RDMA



# CUDA-AWARE MPI

Example:

MPI Rank 0 `MPI_Send` from GPU Buffer

MPI Rank 1 `MPI_Recv` to GPU Buffer

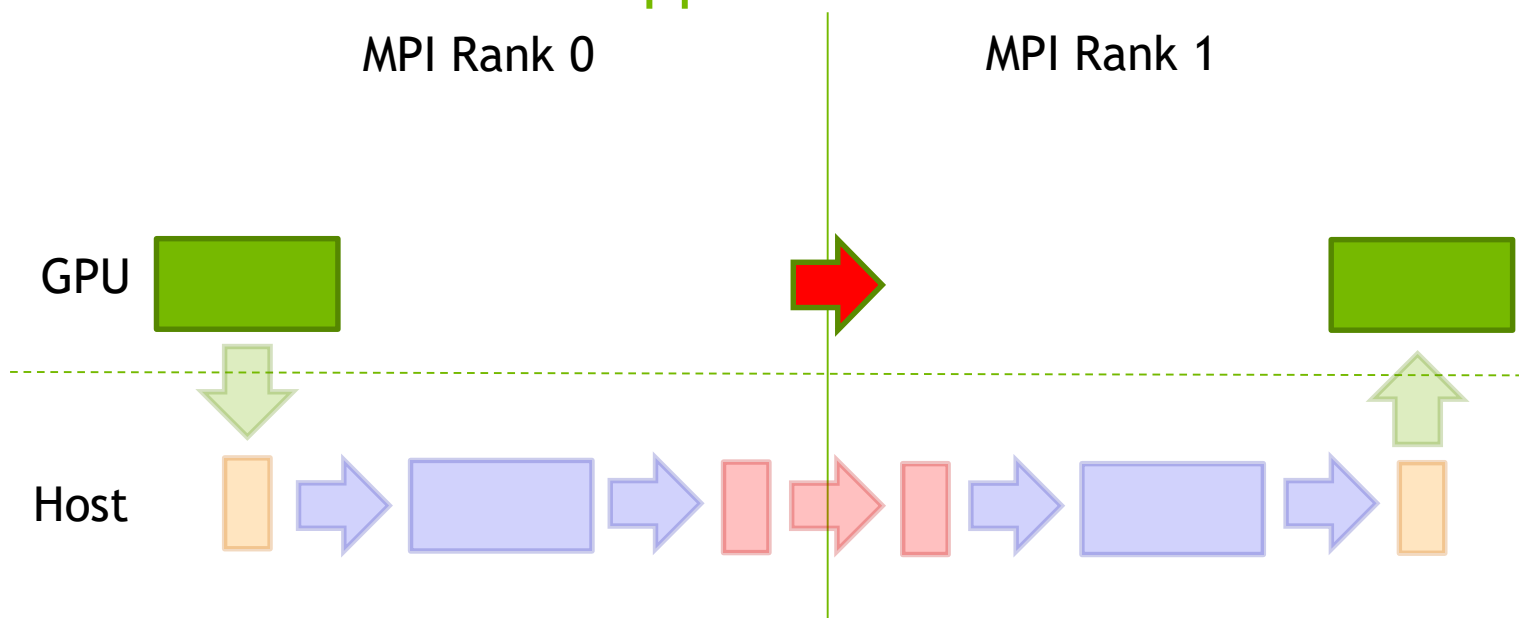
Show how CUDA+MPI works in principle

Depending on the MPI implementation, message size, system setup, ...  
situation might be different

Two GPUs in two nodes

# MPI GPU TO REMOTE GPU

Support for RDMA



```
MPI_Send(s_buf_d,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
```

```
MPI_Recv(r_buf_d,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
```

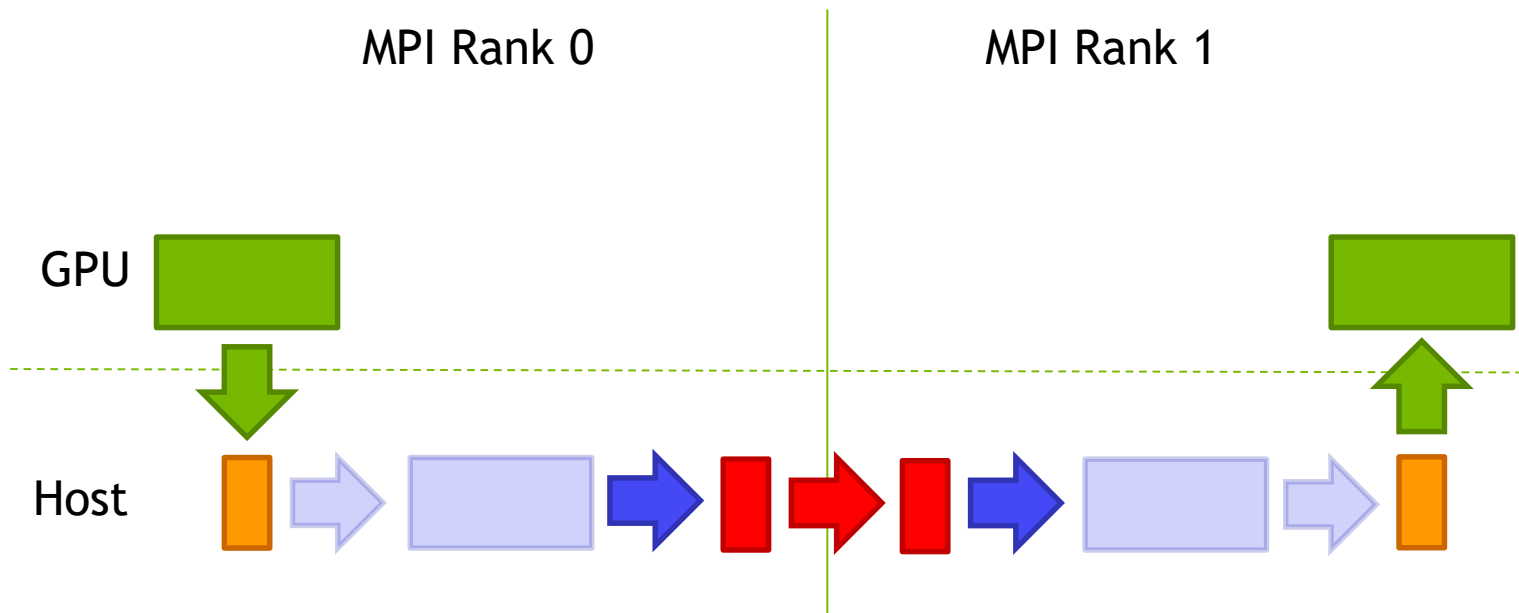
# MPI GPU TO REMOTE GPU

Support for RDMA



# MPI GPU TO REMOTE GPU

without GPUDirect

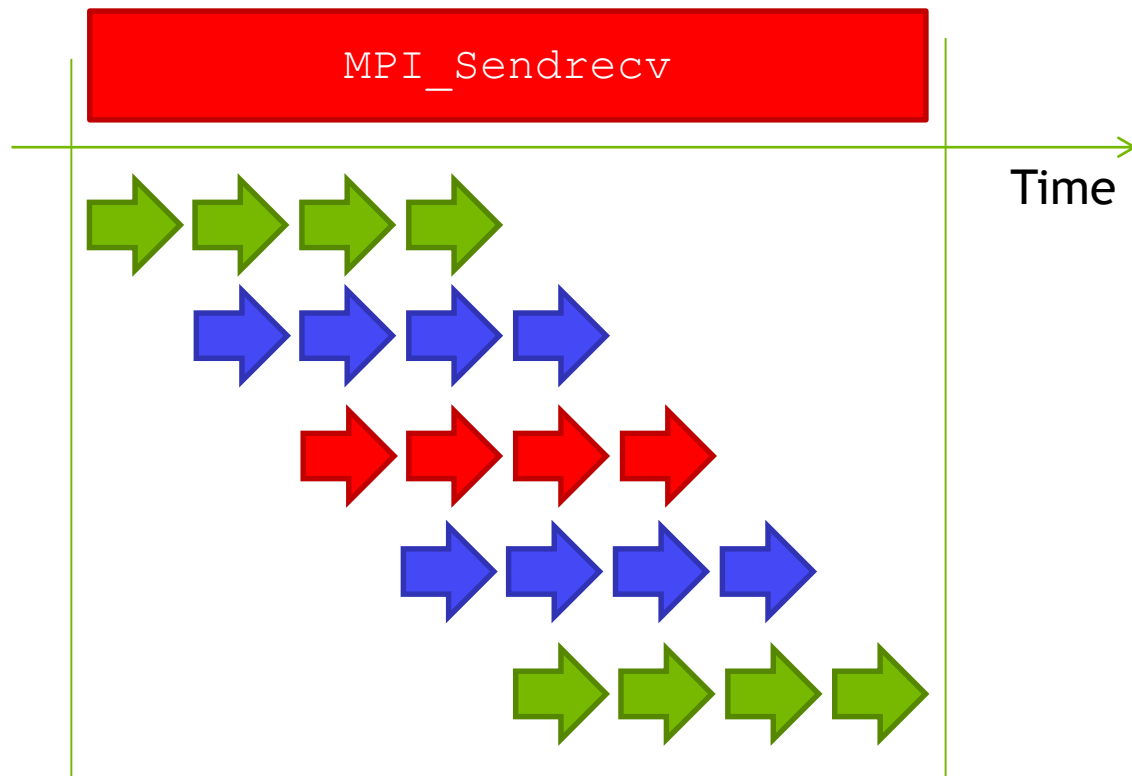


```
MPI_Send(s_buf_d,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
```

```
MPI_Recv(r_buf_d,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
```


# MPI GPU TO REMOTE GPU

without GPUDirect





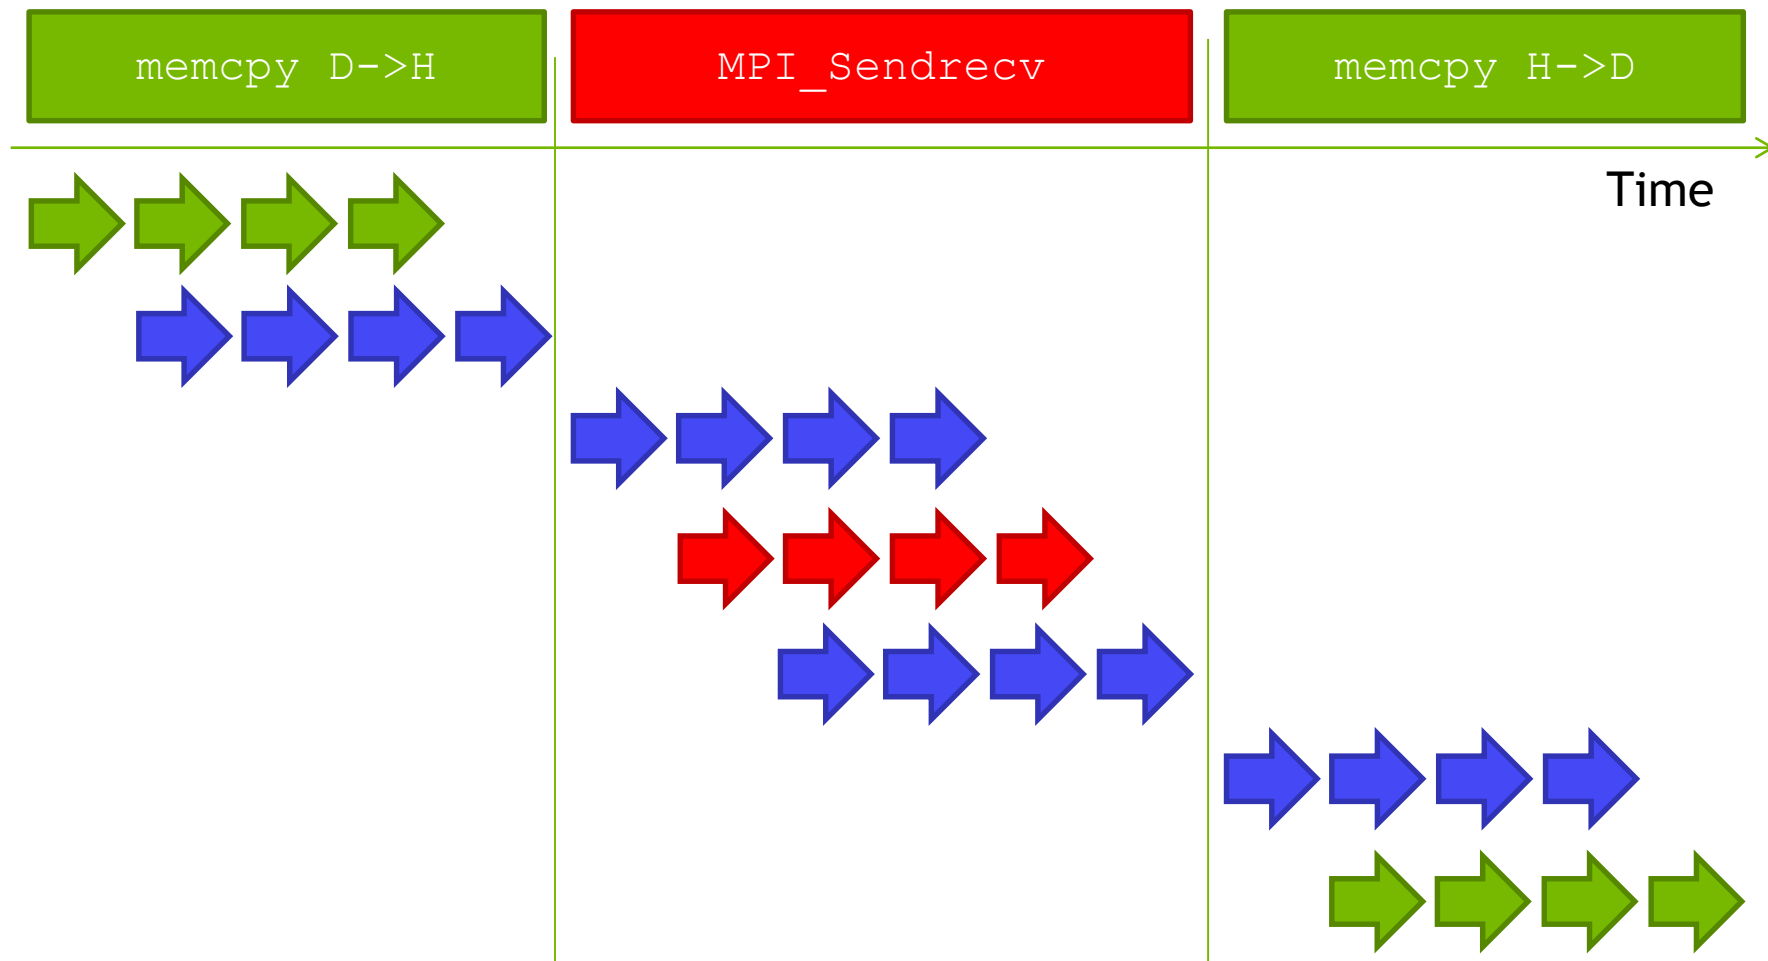
# REGULAR MPI GPU TO REMOTE GPU



```
cudaMemcpy(s_buf_h,s_buf_d,size,cudaMemcpyDeviceToHost);  
MPI_Send(s_buf_h,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
```


```
MPI_Recv(r_buf_h,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);  
cudaMemcpy(r_buf_d,r_buf_h,size,cudaMemcpyHostToDevice);
```

# REGULAR MPI GPU TO REMOTE GPU



# PERFORMANCE RESULTS GPUDIRECT RDMA


MVAPICH2-GDR 2.3a DGX-1V Tesla V100



Latency (1 Byte)      16.75 us      18.68 us      3.25 us

# PERFORMANCE RESULTS GPUDIRECT P2P

MVAPICH2-GDR 2.3a DGX-1V Tesla V100



# MULTI PROCESS SERVICE (MPS) FOR MPI APPLICATIONS

# GPU ACCELERATION OF LEGACY MPI APPS

Typical legacy application

- MPI parallel

- Single or few threads per MPI rank (e.g. OpenMP)

Running with multiple MPI ranks per node

GPU acceleration in phases


- Proof of concept prototype, ...

- Great speedup at kernel level

Application performance misses expectations

# MULTI PROCESS SERVICE (MPS)

For Legacy MPI Applications

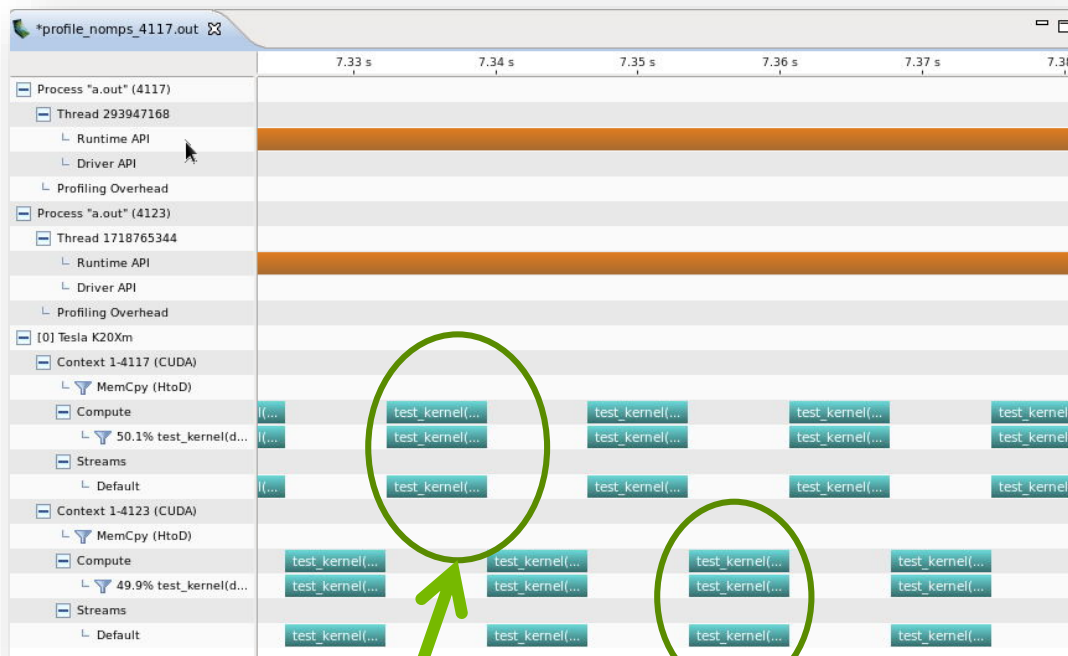



Multicore CPU only

GPU-accelerated

# PROCESSES SHARING GPU WITHOUT MPS

No Overlap



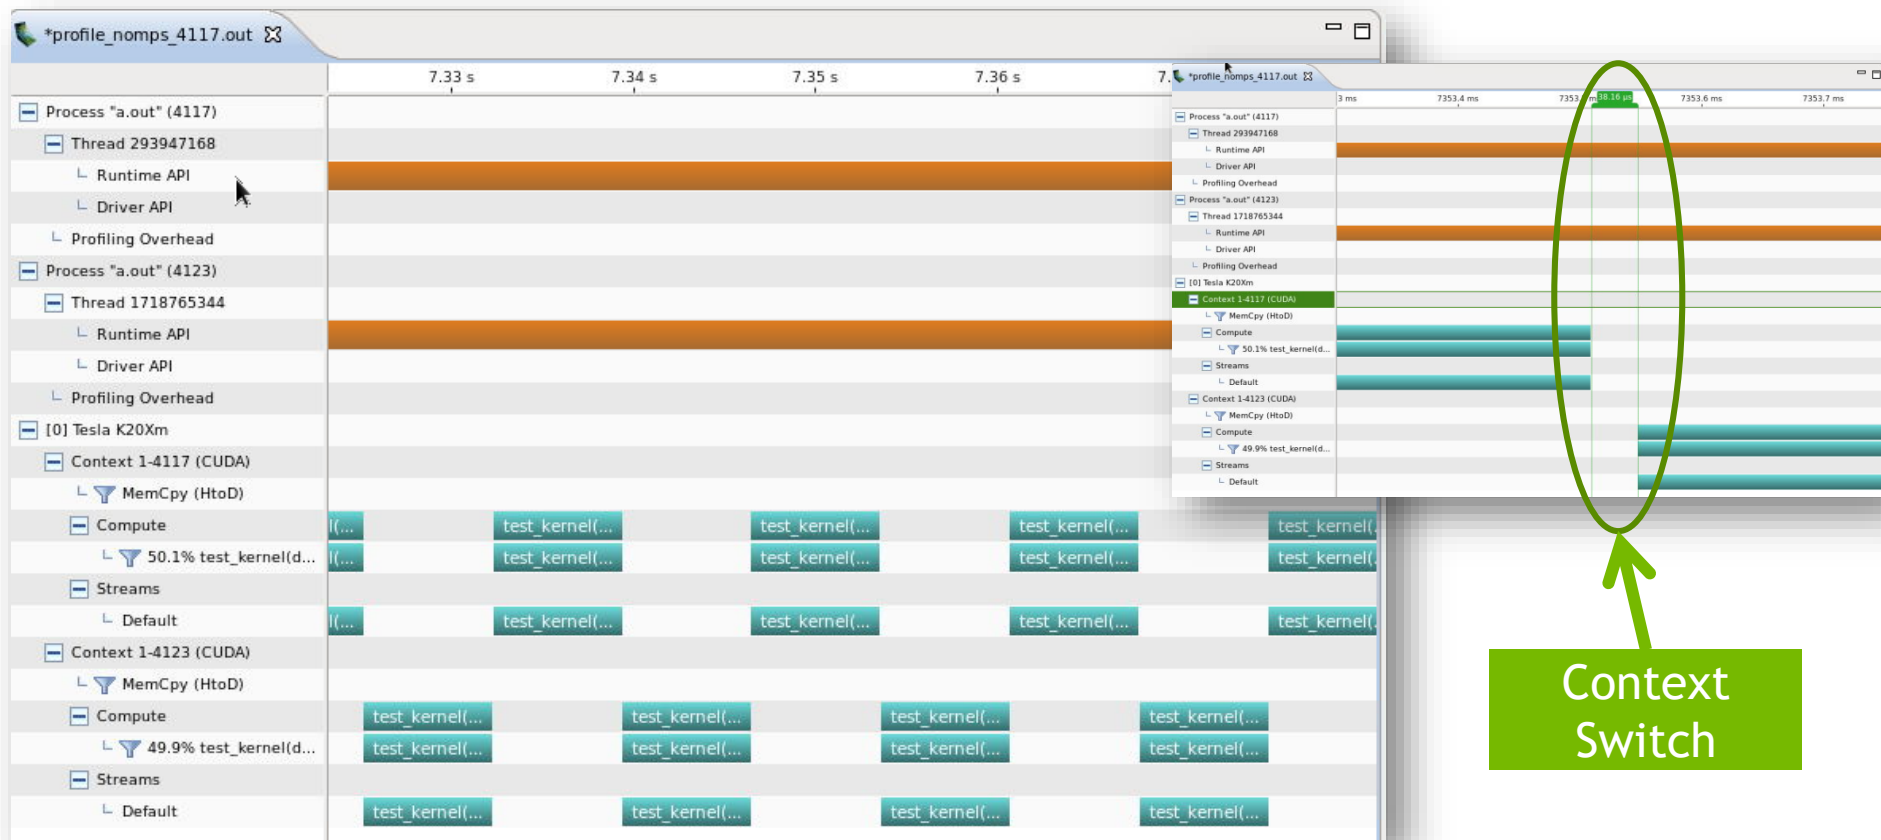
Process A

Process B



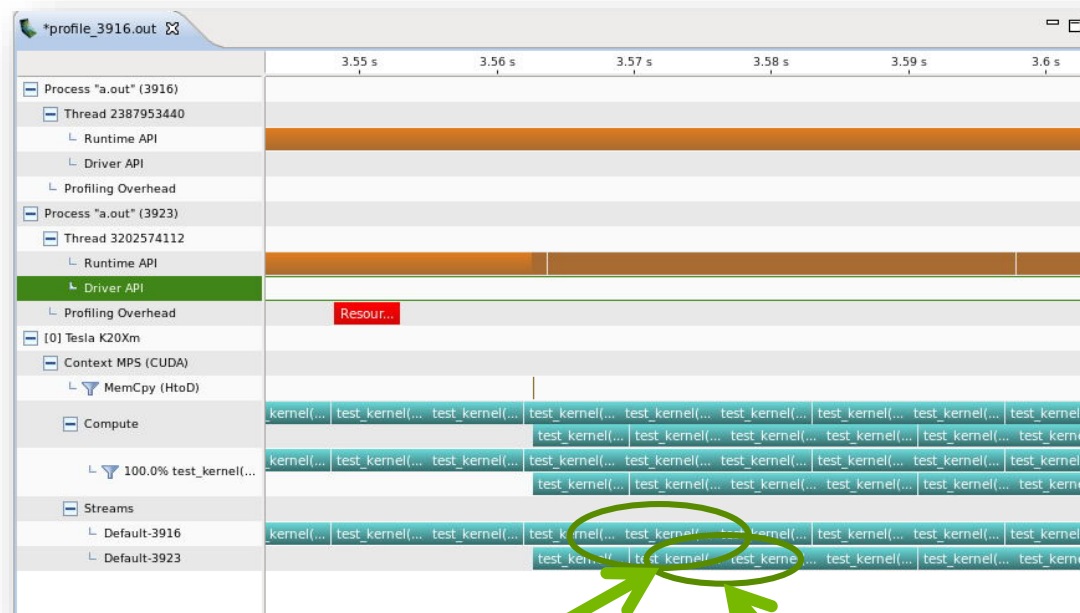
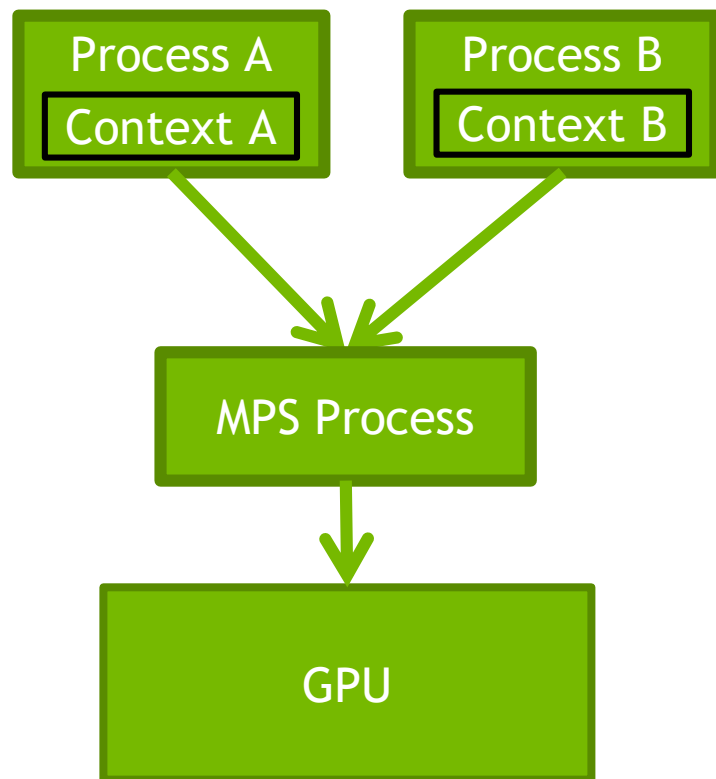
# PROCESSES SHARING GPU WITHOUT MPS

## Context Switch Overhead



# PROCESSES SHARING GPU WITH MPS

## Maximum Overlap

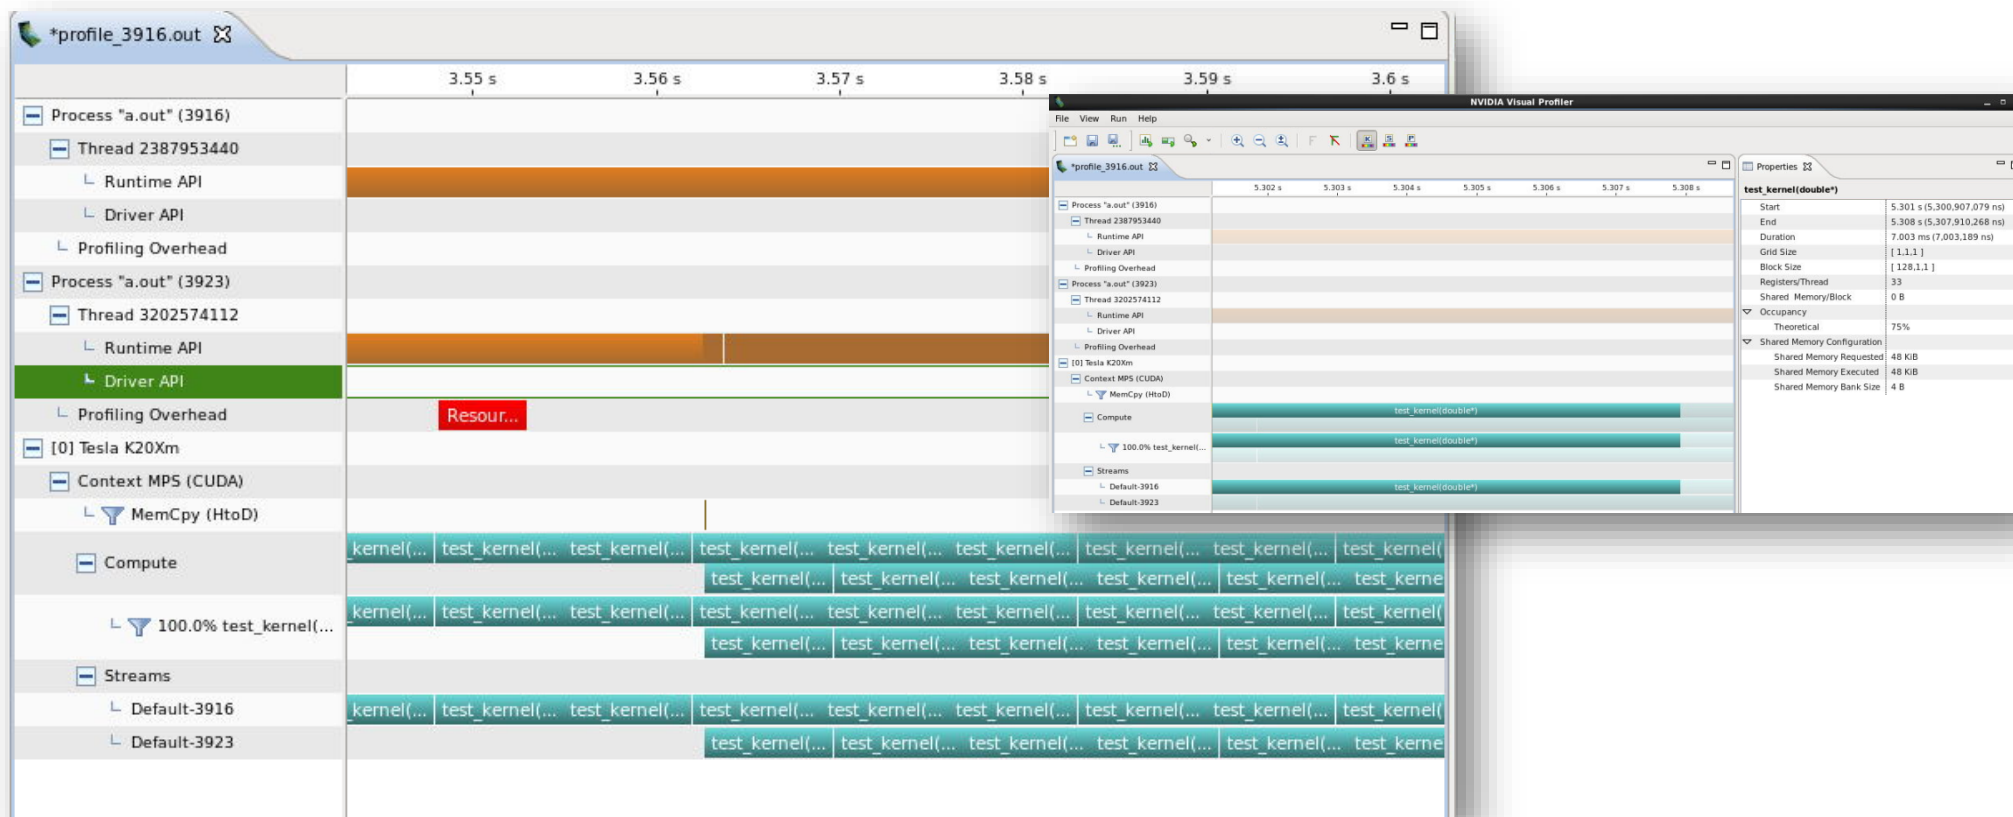


Kernels from  
Process A

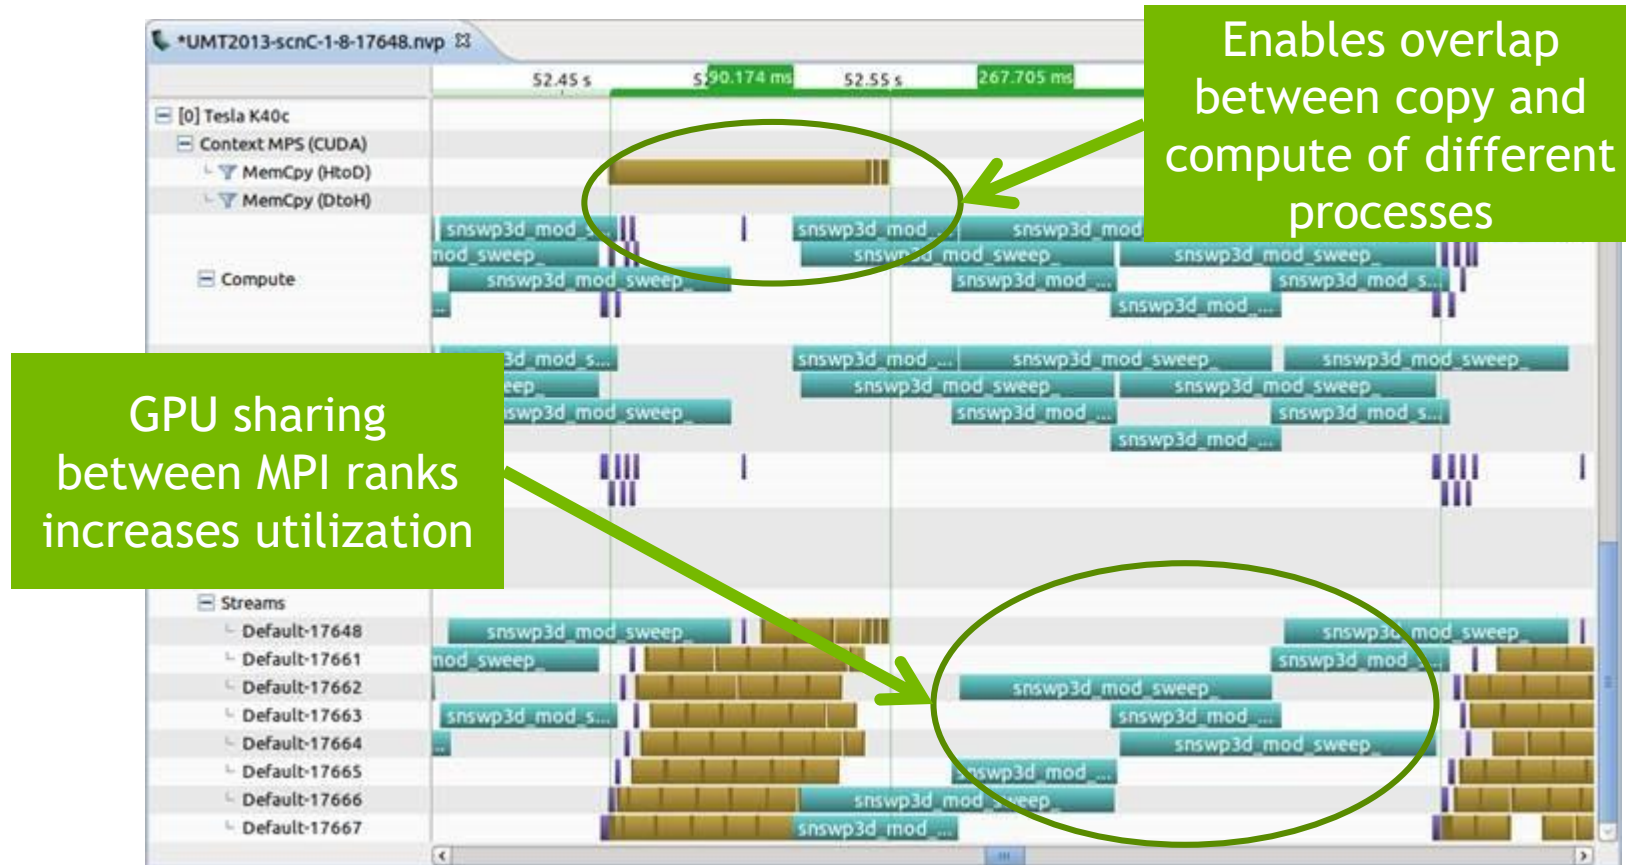
Kernels from  
Process B

# PROCESSES SHARING GPU WITH MPS

## No Context Switch Overhead




# HYPER-Q/MPS CASE STUDY: UMT




# HYPER-Q/MPS CASE STUDIES

## CPU Scaling Speedup



# HYPER-Q/MPS CASE STUDIES

Additional Speedup with MPS



# USING MPS

No application modifications necessary

Not limited to MPI applications

MPS control daemon

Spawn MPS server upon CUDA  
application startup

`#Typical Setup`

```
nvidia-smi -c EXCLUSIVE_PROCESS
```

```
nvidia-cuda-mps-control -d
```

`#On Cray XK/XC systems`

```
export CRAY_CUDA_MPS=1
```

# MPS: IMPROVEMENTS WITH VOLTA

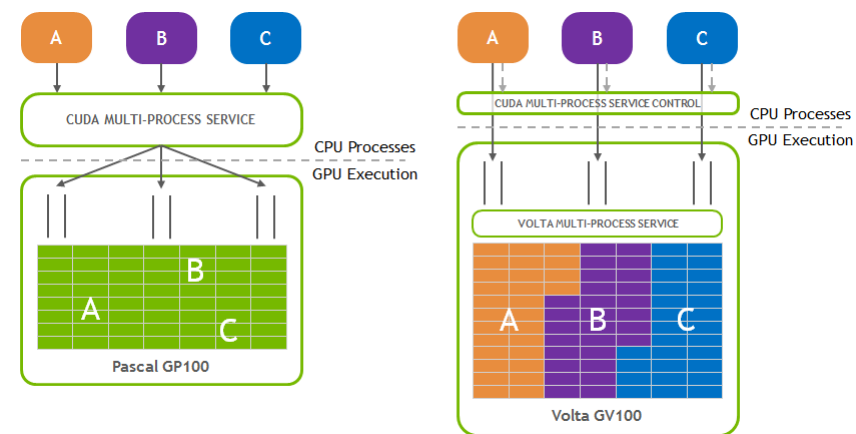
**More MPS clients per GPU:** 48 instead of 16

**Less overhead:** Volta MPS clients submit work directly to the GPU without passing through the MPS server.

**More security:** Each Volta MPS client owns its own GPU address space instead of sharing GPU address space with all other MPS clients.

**More control:** Volta MPS supports limited execution resource provisioning for Quality of Service (QoS). ->

`CUDA_MPS_ACTIVE_THREAD_PERCENTAGE`



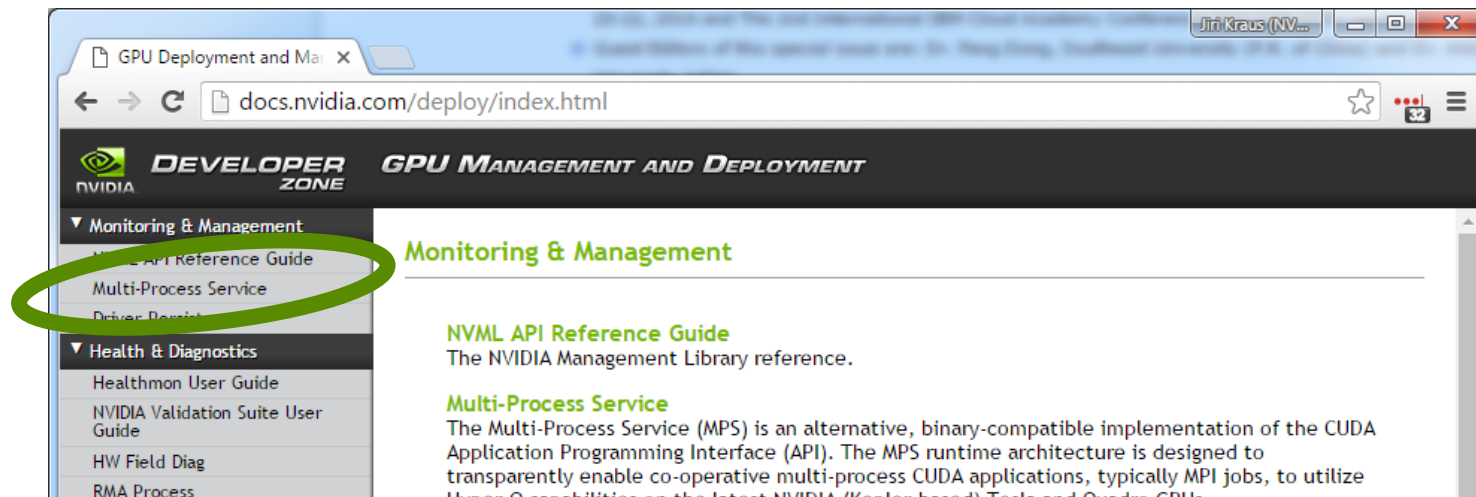


# MPS SUMMARY

Easy path to get GPU acceleration for legacy applications

Enables overlapping of memory copies and compute between different MPI ranks

Remark: MPS adds some overhead!



# DEBUGGING AND PROFILING

# TOOLS FOR MPI+CUDA APPLICATIONS

Memory checking: `cuda-memcheck`

Debugging: `cuda-gdb`

Profiling: `nvprof` and the NVIDIA Visual Profiler (`nvvp`)

# MEMORY CHECKING WITH CUDA-MEMCHECK

cuda-memcheck is a tool similar to Valgrind's memcheck

Can be used in a MPI environment

```
mpiexec -np 2 cuda-memcheck ./myapp <args>
```

Problem: Output of different processes is interleaved

Solution: Use save or log-file command line options

OpenMPI: OMPI\_COMM\_WORLD\_RANK

MPVAPICH2: MV2\_COMM\_WORLD\_RANK

```
mpirun -np 2 cuda-memcheck \
    --log-file name.%q{OMPI_COMM_WORLD_RANK}.log \
    --save name.%q{OMPI_COMM_WORLD_RANK}.memcheck \
    ./myapp <args>
```

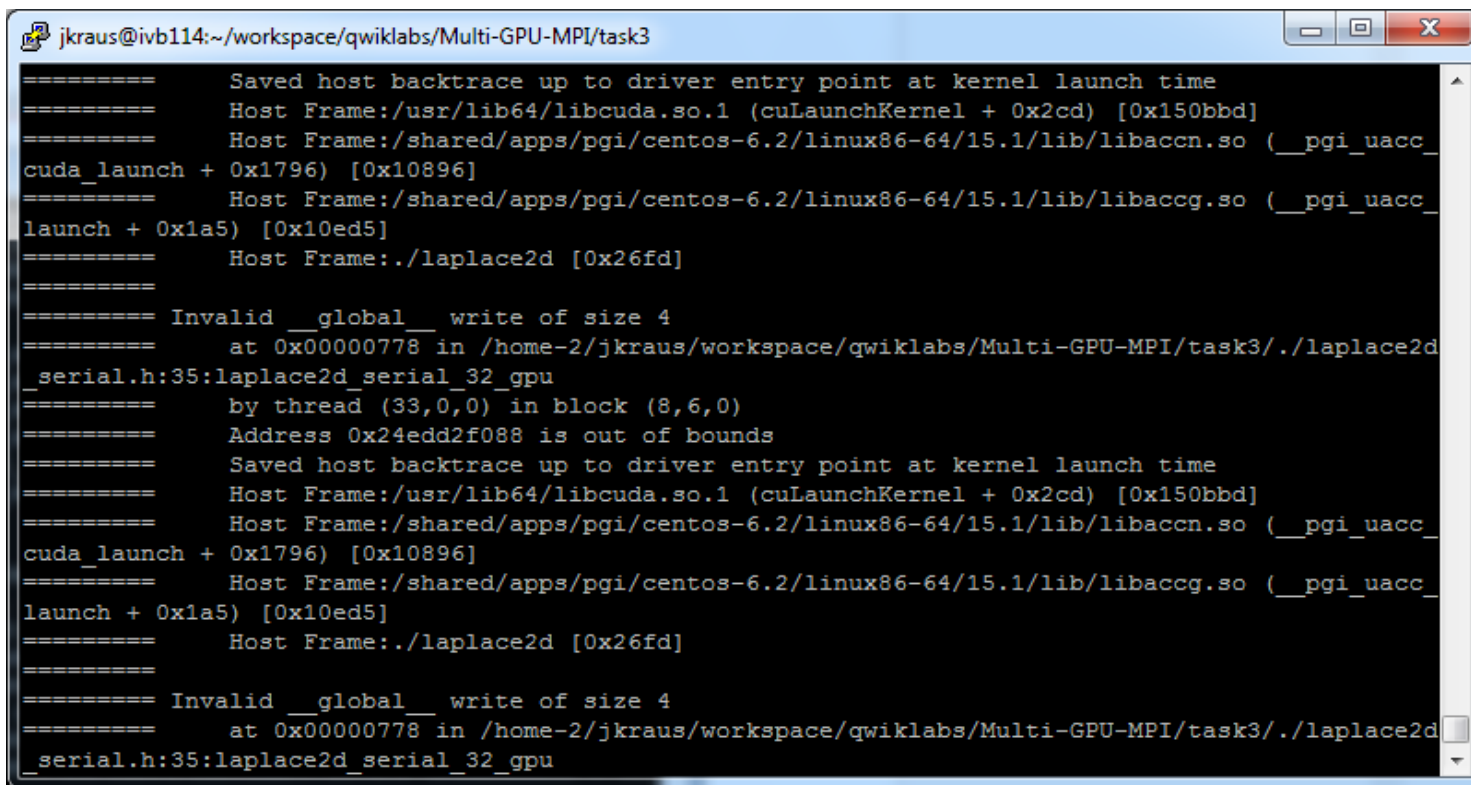
# MEMORY CHECKING WITH CUDA-MEMCHECK

```
jkraus@ivb114:~/workspace/qwiklabs/Multi-GPU-MPI/task3
[jkraus@ivb114 task3]$ mpirun -np 2 cuda-memcheck --log-file laplace2d.%q{OMPI_COMM_WORLD_RANK}.log --save laplace2d.%q{OMPI_COMM_WORLD_RANK}.memcheck ./laplace2d
Jacobi relaxation Calculation: 2048 x 2048 mesh
Calculate reference solution and time serial execution.
call to cuMemcpyDtoHAsync returned error 719: Launch failed (often invalid pointer dereference)
call to cuMemcpyDtoHAsync returned error 719: Launch failed (often invalid pointer dereference)
-----
Primary job terminated normally, but 1 process returned
a non-zero exit code.. Per user-direction, the job has been aborted.
-----
mpirun detected that one or more processes exited with non-zero status, thus causing
the job to be terminated. The first process to do so was:

Process name: [[42894,1],0]
Exit code: 1
-----
[jkraus@ivb114 task3]$ ls laplace2d.*.log laplace2d.*.memcheck
laplace2d.0.log laplace2d.0.memcheck laplace2d.1.log laplace2d.1.memcheck
[jkraus@ivb114 task3]$
```

# MEMORY CHECKING WITH CUDA-MEMCHECK

Read Output Files with `cuda-memcheck --read`

A screenshot of a terminal window with a blue title bar. The title bar text is 'jkraus@ivb114:~/workspace/qwiklabs/Multi-GPU-MPI/task3'. The terminal content shows the output of a CUDA-MEMCHECK run. It includes a host backtrace at kernel launch time, an 'Invalid \_\_global\_\_ write of size 4' error at address 0x24edd2f088, and a second identical backtrace and error message. The error message specifies the location as 'laplace2d\_serial.h:35:laplace2d\_serial\_32\_gpu' and the thread as '(33,0,0)'.

```
===== Saved host backtrace up to driver entry point at kernel launch time
===== Host Frame:/usr/lib64/libcuda.so.1 (cuLaunchKernel + 0x2cd) [0x150bbd]
===== Host Frame:/shared/apps/pgi/centos-6.2/linux86-64/15.1/lib/libaccn.so (__pgi_uacc_
cuda_launch + 0x1796) [0x10896]
===== Host Frame:/shared/apps/pgi/centos-6.2/linux86-64/15.1/lib/libaccg.so (__pgi_uacc_
launch + 0x1a5) [0x10ed5]
===== Host Frame:./laplace2d [0x26fd]
=====
===== Invalid __global__ write of size 4
===== at 0x00000778 in /home-2/jkraus/workspace/qwiklabs/Multi-GPU-MPI/task3/./laplace2d
_serial.h:35:laplace2d_serial_32_gpu
===== by thread (33,0,0) in block (8,6,0)
===== Address 0x24edd2f088 is out of bounds
===== Saved host backtrace up to driver entry point at kernel launch time
===== Host Frame:/usr/lib64/libcuda.so.1 (cuLaunchKernel + 0x2cd) [0x150bbd]
===== Host Frame:/shared/apps/pgi/centos-6.2/linux86-64/15.1/lib/libaccn.so (__pgi_uacc_
cuda_launch + 0x1796) [0x10896]
===== Host Frame:/shared/apps/pgi/centos-6.2/linux86-64/15.1/lib/libaccg.so (__pgi_uacc_
launch + 0x1a5) [0x10ed5]
===== Host Frame:./laplace2d [0x26fd]
=====
===== Invalid __global__ write of size 4
===== at 0x00000778 in /home-2/jkraus/workspace/qwiklabs/Multi-GPU-MPI/task3/./laplace2d
_serial.h:35:laplace2d_serial_32_gpu
```

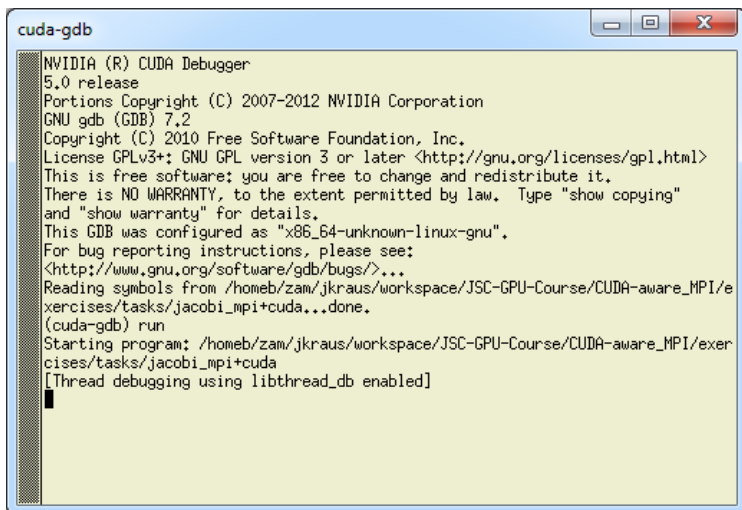
# DEBUGGING MPI+CUDA APPLICATIONS

Using `cuda-gdb` with MPI Applications

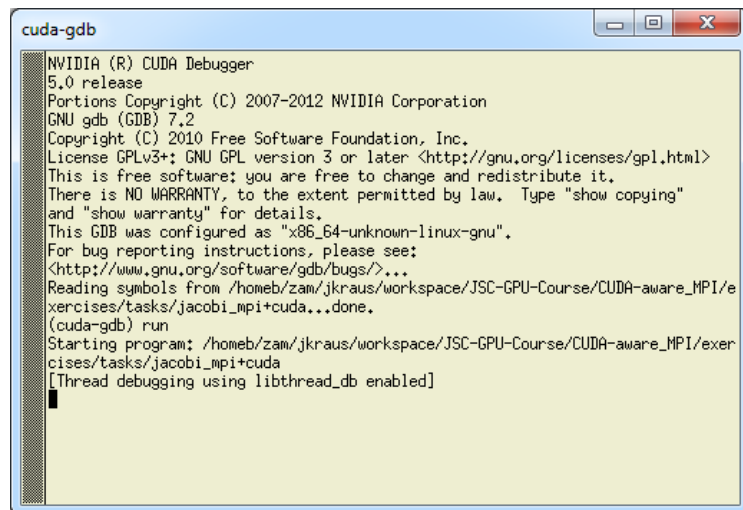
Use `cuda-gdb` just like `gdb`

For smaller applications, just launch `xterms` and `cuda-gdb`

```
mpiexec -x -np 2 xterm -e cuda-gdb ./myapp <args>
```



```
cuda-gdb
NVIDIA (R) CUDA Debugger
5.0 release
Portions Copyright (C) 2007-2012 NVIDIA Corporation
GNU gdb (GDB) 7.2
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software; you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-unknown-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>...
Reading symbols from /homeb/zam/jkraus/workspace/JSC-GPU-Course/CUDA-aware_MPI/e
xercises/tasks/jacobi_mpi+cuda...done.
(cuda-gdb) run
Starting program: /homeb/zam/jkraus/workspace/JSC-GPU-Course/CUDA-aware_MPI/exe
rcises/tasks/jacobi_mpi+cuda
[Thread debugging using libthread_db enabled]
█
```



```
cuda-gdb
NVIDIA (R) CUDA Debugger
5.0 release
Portions Copyright (C) 2007-2012 NVIDIA Corporation
GNU gdb (GDB) 7.2
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software; you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-unknown-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>...
Reading symbols from /homeb/zam/jkraus/workspace/JSC-GPU-Course/CUDA-aware_MPI/e
xercises/tasks/jacobi_mpi+cuda...done.
(cuda-gdb) run
Starting program: /homeb/zam/jkraus/workspace/JSC-GPU-Course/CUDA-aware_MPI/exe
rcises/tasks/jacobi_mpi+cuda
[Thread debugging using libthread_db enabled]
█
```

# DEBUGGING MPI+CUDA APPLICATIONS

## cuda-gdb Attach

```
if ( rank == 0 ) {  
    int i=0;  
    printf("rank %d: pid %d on %s ready for attach\n.", rank, getpid(),name);  
    while (0 == i) { sleep(5); }  
}
```

```
> mpiexec -np 2 ./jacobi_mpi+cuda
```

Jacobi relaxation Calculation: 4096 x 4096 mesh with 2 processes and one Tesla M2070 for each process (2049 rows per process).

rank 0: pid 30034 on judge107 ready for attach

```
> ssh judge107
```

```
jkraus@judge107:~> cuda-gdb --pid 30034
```



# DEBUGGING MPI+CUDA APPLICATIONS

## CUDA\_DEVICE\_WAITS\_ON\_EXCEPTION

```
jkraus@sb077:~/workspace/Jacobi/main/bin
Iteration: 700 - Residue: 0.306564
Iteration: 800 - Residue: 0.306564
Iteration: 900 - Residue: 0.306564
Stopped after 1000 iterations with residue 0.306564
Total Jacobi run time: 0.8700 sec.
Average per-process communication time: 0.2765 sec.
Measured lattice updates: 4.81 GLU/s (total), 1.20 GLU/s (per process)
Measured FLOPS: 24.06 GFLOPS (total), 6.01 GFLOPS (per process)
Measured device bandwidth: 230.95 GB/s (total), 57.74 GB/s (per process)
[jkraus@sb077 bin]$ CUDA_DEVICE_WAITS_ON_EXCEPTION=1 MV2_USE_
aware_mpi_async -t 2 2 -d 1024 1024 -fs
Topology size: 2 x 2
Local domain size (current node): 1024 x 1024
Global domain size (all nodes): 2048 x 2048
Starting Jacobi run with 4 processes:
sb077: The application encountered a device error and CUDA DE
can now attach a debugger to the application (PID 28250) for
sb077: The application encountered a device error and CUDA DE
can now attach a debugger to the application (PID 28252) for
sb077: The application encountered a device error and CUDA DE
can now attach a debugger to the application (PID 28251) for
sb077: The application encountered a device error and CUDA DE
can now attach a debugger to the application (PID 28249) for

Reading symbols from /usr/lib64/libnes-rdmav2.so...(no debugging symbols found)...done.
Loaded symbols for /usr/lib64/libnes-rdmav2.so
Reading symbols from /usr/lib64/libmlx4-rdmav2.so...(no debugging symbols found)...done.
Loaded symbols for /usr/lib64/libmlx4-rdmav2.so
Reading symbols from /usr/lib64/libipathverbs-rdmav2.so...(no debugging symbols found)...done.
Loaded symbols for /usr/lib64/libipathverbs-rdmav2.so
0x00007f5ba011fa01 in clock_gettime ()
$1 = 1
CUDA Exception: Device Illegal Address
The exception was triggered in device 3.
Program received signal CUDA_EXCEPTION_10, Device Illegal Address.
[Switching focus to CUDA kernel 0, grid 8, block (6,36,0), thread (0,6,0), device 3, sm 0, warp
13, lane 0]
0x00000000018e1ce8 in JacobiComputeKernel<<<(64,64,1),(16,16,1)>>> (size=..., startmod=...,
endmod=..., oldBlock=0x2300200000, newBlock=0x2300b20000, devResidue=0x2301340000,
stride=1024) at Device.cu:150
150 AtomicMax<real>(devResidue, rabs(newVal - oldBlock[memIdx]));
(cuda-gdb) bt
#0 0x00000000018e1ce8 in JacobiComputeKernel<<<(64,64,1),(16,16,1)>>> (size=...,
startmod=..., endmod=..., oldBlock=0x2300200000, newBlock=0x2300b20000,
devResidue=0x2301340000, stride=1024) at Device.cu:150
(cuda-gdb)
```

# DEBUGGING MPI+CUDA APPLICATIONS

With `CUDA_ENABLE_COREDUMP_ON_EXCEPTION=1` core dumps are generated in case of an exception:

- Can be used for offline debugging

- Helpful if live debugging is not possible

`CUDA_ENABLE_CPU_COREDUMP_ON_EXCEPTION`: Enable/Disable CPU part of core dump (enabled by default)

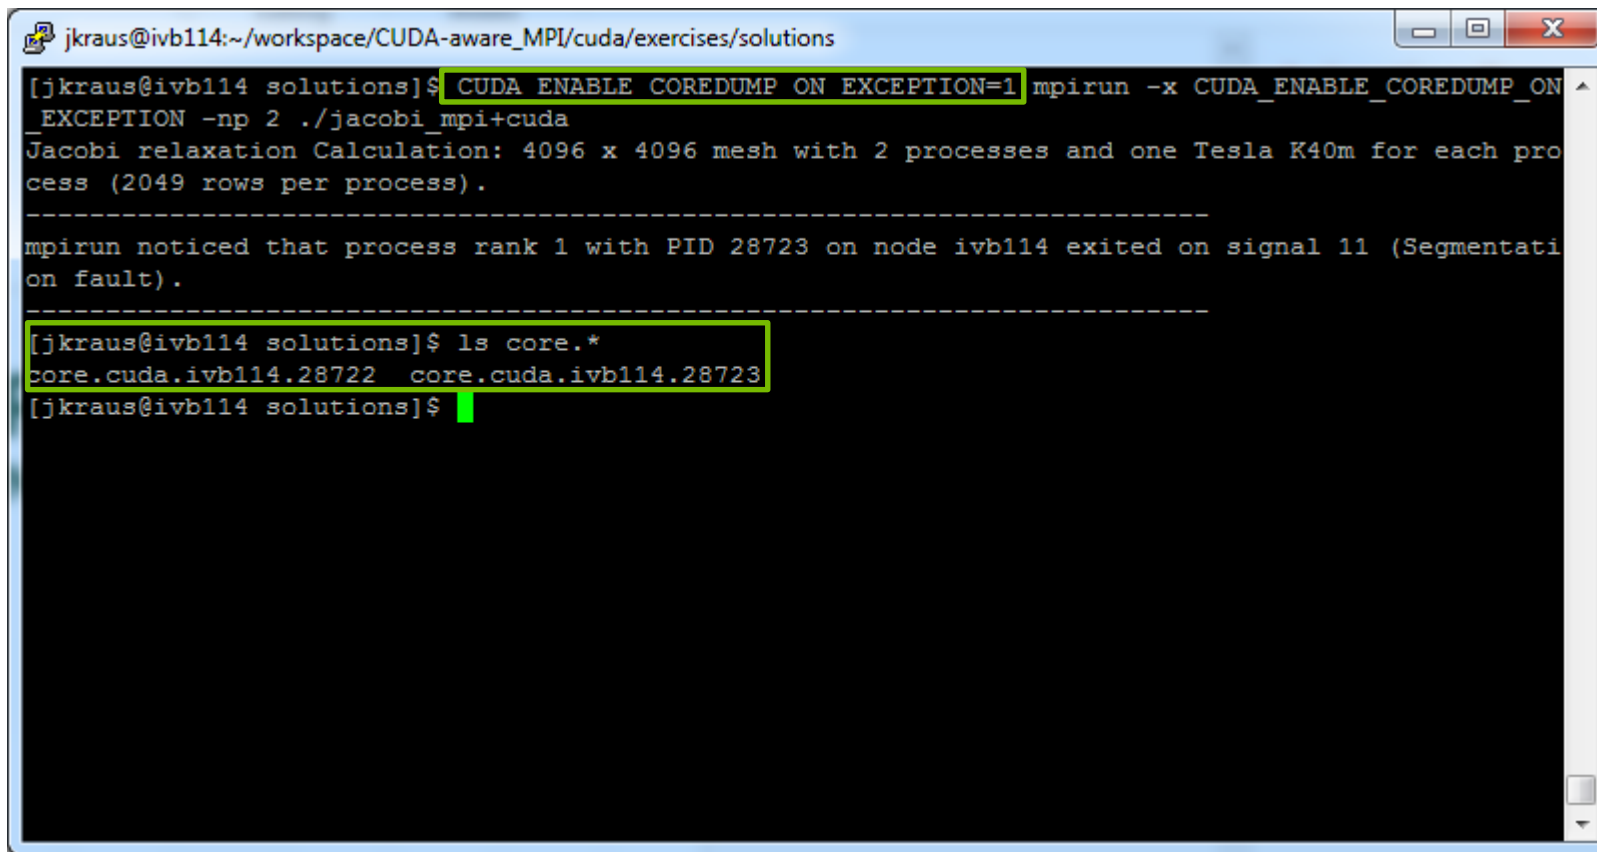
`CUDA_COREDUMP_FILE`: Specify name of core dump file

Open GPU: `(cuda-gdb) target cudacore core.cuda`

Open CPU+GPU: `(cuda-gdb) target core core.cpu core.cuda`

# DEBUGGING MPI+CUDA APPLICATIONS

CUDA\_ENABLE\_COREDUMP\_ON\_EXCEPTION



```
jkraus@ivb114:~/workspace/CUDA-aware_MPI/cuda/exercises/solutions
[jkraus@ivb114 solutions]$ CUDA_ENABLE_COREDUMP_ON_EXCEPTION=1 mpirun -x CUDA_ENABLE_COREDUMP_ON_EXCEPTION -np 2 ./jacobi_mpi+cuda
Jacobi relaxation Calculation: 4096 x 4096 mesh with 2 processes and one Tesla K40m for each process (2049 rows per process).

-----
mpirun noticed that process rank 1 with PID 28723 on node ivb114 exited on signal 11 (Segmentation fault).

-----
[jkraus@ivb114 solutions]$ ls core.*
core.cuda.ivb114.28722 core.cuda.ivb114.28723
[jkraus@ivb114 solutions]$
```

# DEBUGGING MPI+CUDA APPLICATIONS

## CUDA\_ENABLE\_COREDUMP\_ON\_EXCEPTION

```
jkraus@ivb114:~/workspace/CUDA-aware_MPI/cuda/exercises/solutions
NVIDIA (R) CUDA Debugger
7.0 release
Portions Copyright (C) 2007-2014 NVIDIA Corporation
GNU gdb (GDB) 7.6.2
Copyright (C) 2013 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-unknown-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
(cuda-gdb) target cudacore core.cuda.ivb114.28722
Opening GPU coredump: core.cuda.ivb114.28722
[New Thread 28742]

CUDA Exception: Device Illegal Address
The exception was triggered in device 0.
[Current focus set to CUDA kernel 0, grid 1, block (107,0,0), thread (0,12,0), device 0, sm 12,
warp 6, lane 0]
#0  0x0000000001c02ac0 in jacobi_kernel<<<(257,129,1),(16,16,1)>>> (u_d=0x23048a0000,
    unew_d=0x23068c0000, n=2049, m=4096, residue_d=0x23088e0000) at jacobi_cuda_kernel.cu:43
43      residue = fabsf(unew_d[j *m+ i]-u_d[j *m+ i]);
(cuda-gdb) █
```


# DEBUGGING MPI+CUDA APPLICATIONS

## Third Party Tools

Allinea DDT debugger

Rogue Wave TotalView

Stacks		
Threads	CUDA Threads	Function
1	0	main (prefix.cu:193)
1	0	cudasummer (prefix.cu:143)
1	0	prefixsum (prefix.cu:105)
1	512	zarro (prefix.cu:89)
1	480	zarro (prefix.cu:90)



# PROFILING MPI+CUDA APPLICATIONS

Using `nvprof`+**NVVP**

New since CUDA 9

Embed MPI rank in output filename, process name, and context name (OpenMPI)

```
mpirun -np $np nvprof --output-profile profile.%q{OMPI_COMM_WORLD_RANK} \
      --process-name "rank %q{OMPI_COMM_WORLD_RANK}" \
      --context-name "rank %q{OMPI_COMM_WORLD_RANK}" \
      --annotate-mpi openmpi
```

Alternatives:

Only save the textual output (`--log-file`)

Collect data from all processes that run on a node (`--profile-all-processes`)

**MPICH2:** `MV2_COMM_WORLD_RANK`  
`--annotate-mpi mpich`

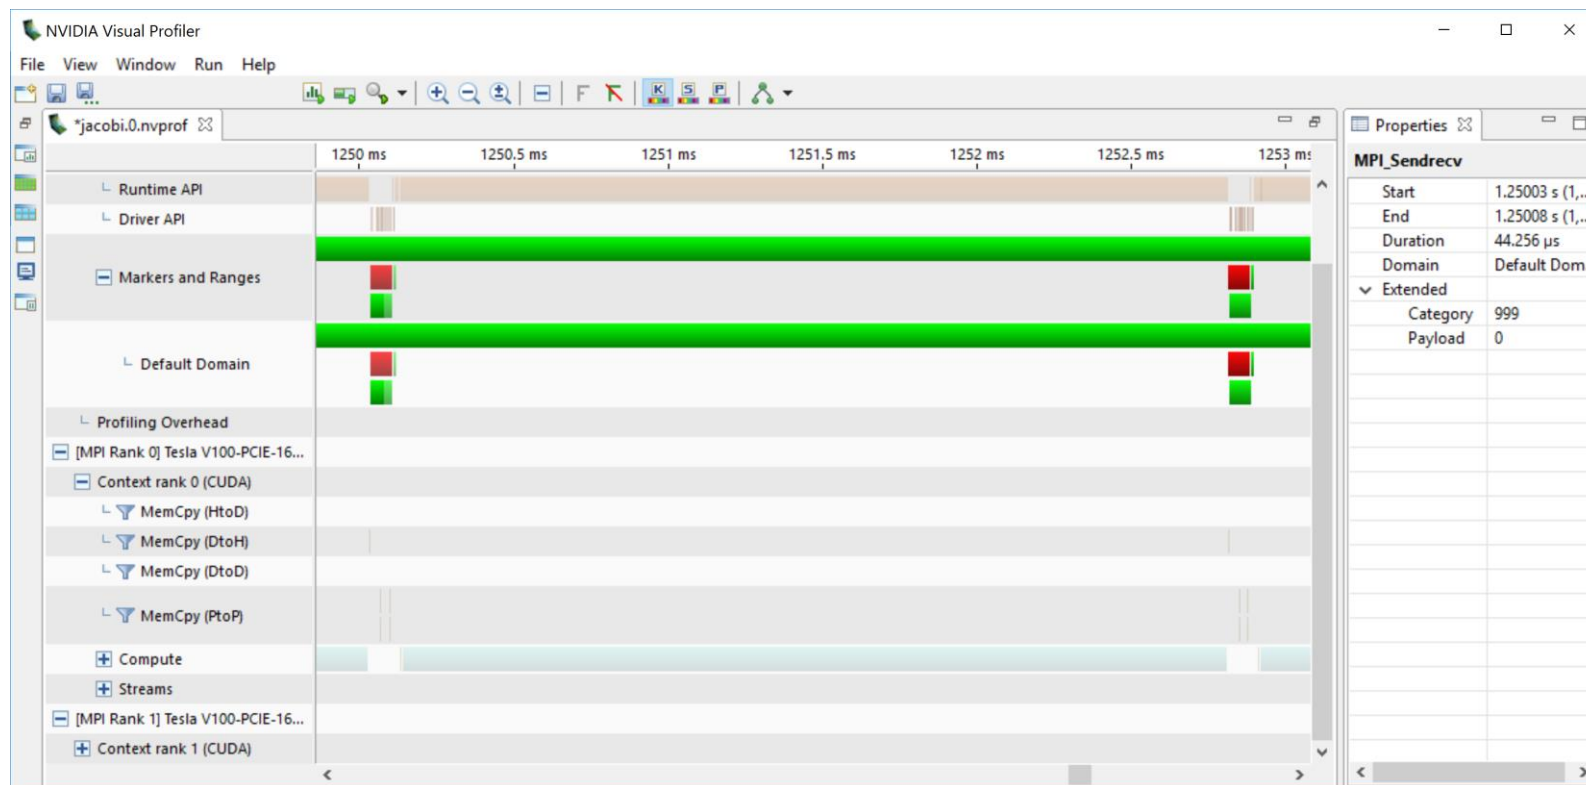
# PROFILING MPI+CUDA APPLICATIONS

Using nvprof+NVVP

```
jakraus@psgcluster:~/workspace/GTC/GTCUS2017/Multi_GPU_Programming_Models/mpi
[jakraus@hsw215 mpi]$ NP=2 make profile
mpirun -np 2 nvprof -o jacobi.%q{OMPI_COMM_WORLD_RANK}.nvprof --context-name "rank %q{OMPI_COMM_WORLD_RANK}" --niter 10 -nx 8192 -ny 8192
==34257== NVPROF is profiling process 34257, command: ./jacobi -niter 10 -nx 8192 -ny 8192
==34256== NVPROF is profiling process 34256, command: ./jacobi -niter 10 -nx 8192 -ny 8192
Single GPU jacobi relaxation: 10 iterations on 8192 x 8192 mesh with norm check every 1 iterations
0, 22.626005
Jacobi relaxation: 10 iterations on 8192 x 8192 mesh with norm check every 1 iterations
0, 22.626020
Num GPUs: 2.
8192x8192: 1 GPU: 0.0600 s. 2 GPUs: 0.0287 s. speedup: 2.09. efficiency: 104.63
==34257== Generated result file: /home/jakraus/workspace/GTC/GTCUS2017/Multi_GPU_Programming_Models/mpi/jacobi.0.nvprof
==34256== Generated result file: /home/jakraus/workspace/GTC/GTCUS2017/Multi_GPU_Programming_Models/mpi/jacobi.1.nvprof
[jakraus@hsw215 mpi]$
```

# PROFILING MPI+CUDA APPLICATIONS

Using `nvprof`+`NVVP`





# PROFILING MPI+CUDA APPLICATIONS

## Third Party Tools

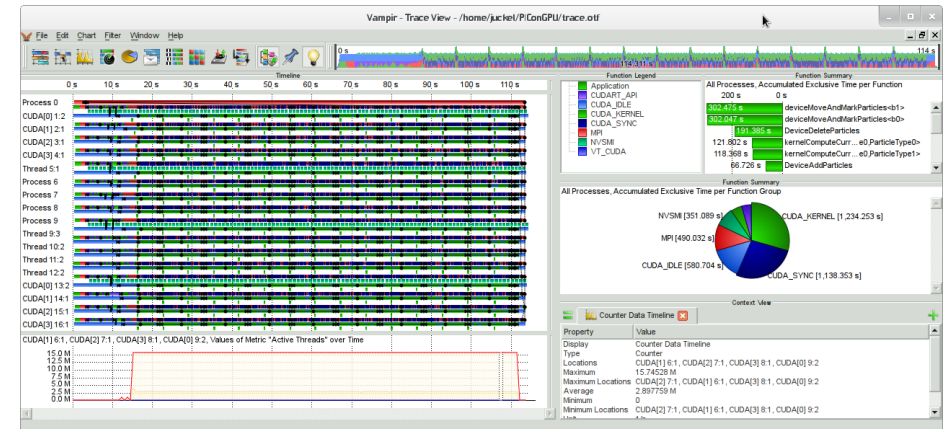

Multiple parallel profiling tools are CUDA-aware

Score-P

Vampir

Tau

These tools are good for discovering MPI issues as well as basic CUDA performance inhibitors.



# ADVANCED MPI ON GPUS

# BEST PRACTICE: USE NON-BLOCKING MPI

BLOCKING

```
#pragma acc host_data use_device ( u_new ) {  
MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,  
             u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,  
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
MPI_Sendrecv(u_new+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,  
             u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,  
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```


NON-BLOCKING

```
MPI_Request t_b_req[4];  
#pragma acc host_data use_device ( u_new ) {  
    MPI_Irecv(u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_b_req+1);  
    MPI_Irecv(u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, t_b_req+2);  
    MPI_Isend(u_new+offset_last_row, m-2, MPI_DOUBLE, t_b_req+3);  
    MPI_Isend(u_new+offset_first_row, m-2, MPI_DOUBLE, t_b_req+4);  
}  
MPI_Waitall(4, t_b_req, MPI_STATUSES_IGNORE);
```

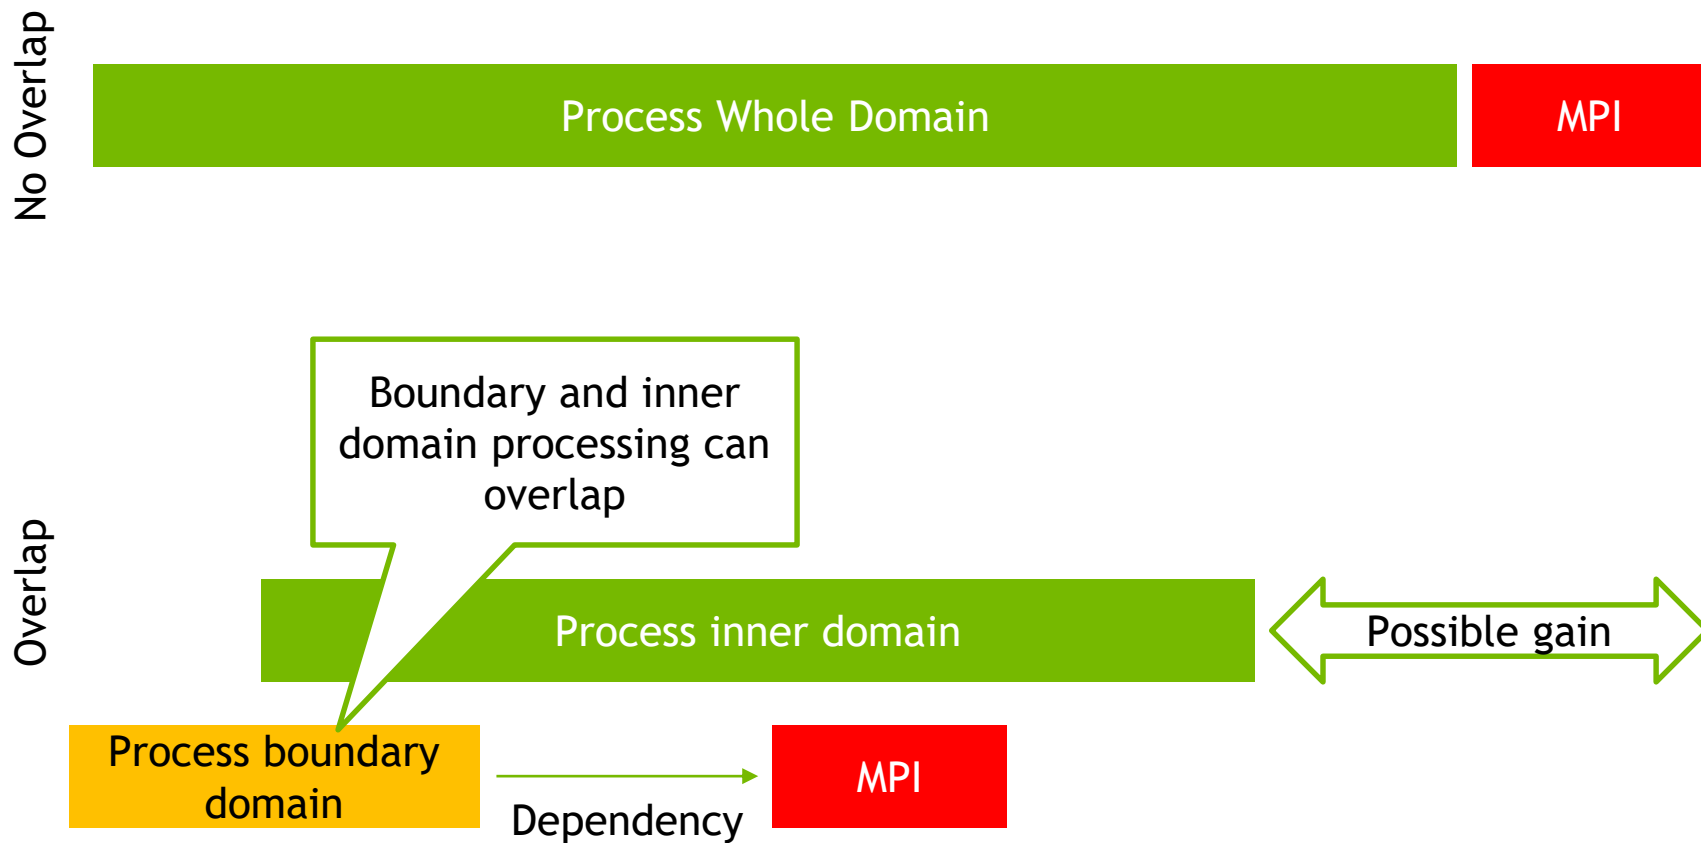
Gives MPI more  
opportunities to build  
efficient pipelines

# COMMUNICATION + COMPUTATION OVERLAP

OpenMPI 3.0.1RC1 - DGX-1V - 2048x2048



# COMMUNICATION + COMPUTATION OVERLAP



# COMMUNICATION + COMPUTATION OVERLAP

## CUDA with Streams

```
process_boundary_and_pack<<<gs_b,bs_b,0,s1>>>(u_new_d,u_d,to_left_d,to_right_d,n,m);

process_inner_domain<<<gs_id,bs_id,0,s2>>>(u_new_d, u_d,to_left_d,to_right_d,n,m);

cudaStreamSynchronize(s1);           //wait for boundary
MPI_Request req[8];

//Exchange halo with left, right, top and bottom neighbor

MPI_Waitall(8, req, MPI_STATUSES_IGNORE);
unpack<<<gs_s,bs_s,0,s2>>>(u_new_d, from_left_d, from_right_d, n, m);

cudaDeviceSynchronize();             //wait for iteration to finish
```


# COMMUNICATION + COMPUTATION OVERLAP

## OpenACC with Async Queues

```
#pragma acc parallel loop present ( u_new, u, to_left, to_right ) async(1)
for ( ... )
    //Process boundary and pack to_left and to_right
#pragma acc parallel loop present ( u_new, u ) async(2)
for ( ... )
    //Process inner domain
#pragma acc wait(1)                //wait for boundary
MPI_Request req[8];
#pragma acc host_data use_device ( from_left, to_left, from_right, to_right, u_new ) {
    //Exchange halo with left, right, top and bottom neighbor
}
MPI_Waitall(8, req, MPI_STATUSES_IGNORE);
#pragma acc parallel loop present ( u_new, from_left, from_right ) async(2)
for ( ... )
    //unpack from_left and from_right
#pragma acc wait                //wait for iteration to finish
```

# COMMUNICATION + COMPUTATION OVERLAP

OpenMPI 3.0.1RC1 - DGX-1V - 2048x2048



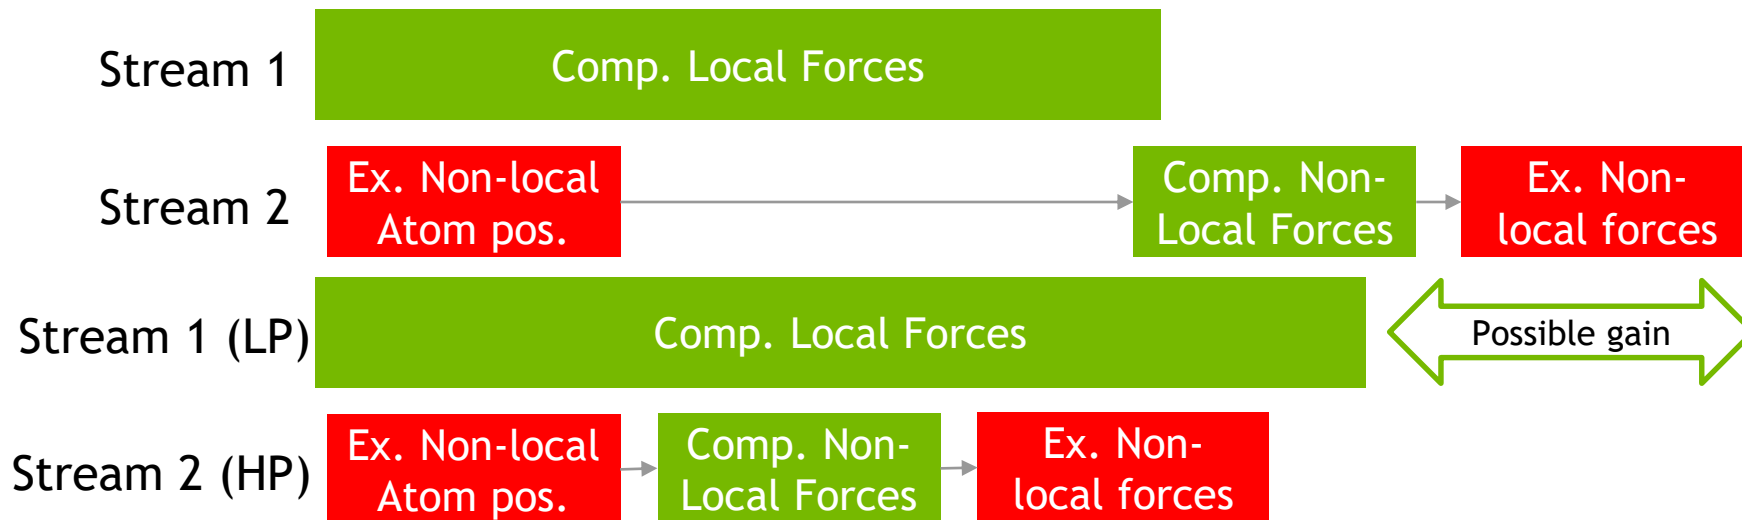


# HIGH PRIORITY STREAMS

Improve scalability with high priority streams

```
__host__ cudaError_t cudaStreamCreateWithPriority ( cudaStream_t* pStream, unsigned  
int flags, int priority )
```

Use-case: MD Simulations



# MPI AND UNIFIED MEMORY

## CAVEAT

Using Unified Memory with a non Unified Memory-aware MPI might fail with errors or even worse silently produce wrong results, e.g. when registering Unified Memory for RDMA.



**Use a Unified Memory-aware MPI,  
e.g. OpenMPI since 1.8.5 or MVAPICH2-GDR since 2.2b**

Unified Memory-aware: CUDA-aware MPI with support for Unified Memory

# MPI AND UNIFIED MEMORY

## Performance Implications

Unified Memory can be used by any processor in the system

Memory pages of a Unified Memory allocation may migrate between processors memories to ensure coherence and maximize performance

Different data paths are optimal for performance depending on where the data is:  
e.g. NVLink between peer GPUs



The MPI implementation needs to know where the data is,  
but it can't!

# MPI AND UNIFIED MEMORY

## Performance Implications - Simple Example

```
cudaMallocManaged( &array, n*sizeof(double), cudaMemAttachGlobal );

while( ... ) {

    foo(array,n);

    MPI_Send(array,...);

    foo(array,n);

}
```

# MPI AND UNIFIED MEMORY

## Performance Implications - Simple Example

- ▶ If foo is a CPU function pages of array might migrate to System Memory
- ▶ If foo is a GPU function pages of array might migrate to GPU Memory
- ▶ The MPI implementation is not aware of the application and thus doesn't know where array is and what's optimal


```
while( ... ) {  
    foo(array,n);  
  
    MPI_Send(array,...);  
  
    foo(array,n);  
}
```

# MPI AND UNIFIED MEMORY

## The Future with Data Usage Hints

Tell where the application intends to use the data

```
cudaMallocManaged( &array, n*sizeof(double), cudaMemAttachGlobal );  
  
cudaMemAdvise(array, n*sizeof(double), cudaMemAdviseSetPreferredLocation, device);  
  
while( ... ) {  
    foo(array, n);  
    MPI_Send(array, ...);  
    foo(array, n);  
}
```



Array is intended to  
be used on the GPU  
with the id device

Remark: Data Usage Hints are available since CUDA 8, but currently not evaluated by any Unified Memory-aware MPI implementation.

# MPI AND UNIFIED MEMORY

## The Future with Data Usage Hints

Tell where the application intends to use the data

```
cudaMallocManaged( &array, n*sizeof(double), cudaMemAttachGlobal );  
  
cudaMemAdvise(array, n*sizeof(double), cudaMemAdviseSetPreferredLocation, cudaCpuDeviceId);  
  
while( ... ) {  
    foo(array, n);  
  
    MPI_Send(array, ...);  
  
    foo(array, n);  
  
}
```



Array is intended to be used on the CPU

Remark: Data Usage Hints are available since CUDA 8, but currently not evaluated by any Unified Memory-aware MPI implementation.

# MPI AND UNIFIED MEMORY

## The Future with Data Usage Hints - Summary

Data usage hints can be queried by the MPI Implementation and allow it to take the optimal data path

If the application lies about the data usage hints it will run correctly but performance will be affected

Performance tools help to identify missing or wrong data usage hints

Data usage hints are general useful for the Unified Memory system and can improve application performance.

Remark: Data Usage Hints are only hints to guide the data usage policies of the Unified Memory system. The Unified Memory system might ignore them, e.g. to ensure coherence or in oversubscription scenarios.



# MPI AND UNIFIED MEMORY

## Current Status

Available Unified Memory-aware MPI implementations

- OpenMPI (since 1.8.5)
- MVAPICH2-GDR (since 2.2b)
  - Performance improvements with 2.2RC1 for Intranode GPU to GPU communication

Currently both don't evaluate Data Usage Hints, i.e. all Unified Memory is treated as Device Memory



Good performance if all buffers used in MPI are touched mainly on the GPU.

# MPI AND UNIFIED MEMORY

## Without Unified Memory-aware MPI

Only use non Unified Memory Buffers for MPI: `cudaMalloc`, `cudaMallocHost` or `malloc`

Application managed non Unified Memory Buffers also allow to work around current missing cases in Unified Memory-aware MPI Implementations.

# DETECTING CUDA-AWARENESS

OpenMPI (since 2.0.0):

Macro:

```
MPIX_CUDA_AWARE_SUPPORT
```

Function for runtime decisions

```
MPIX_Query_cuda_support()
```

Include `mpi-ext.h` for both.

See <http://www.open-mpi.org/faq/?category=runcuda#mpi-cuda-aware-support>

